

# Natural Language Processing

---

**L2 / 17.02.26**

- [Course page](#)
- [GitHub](#)

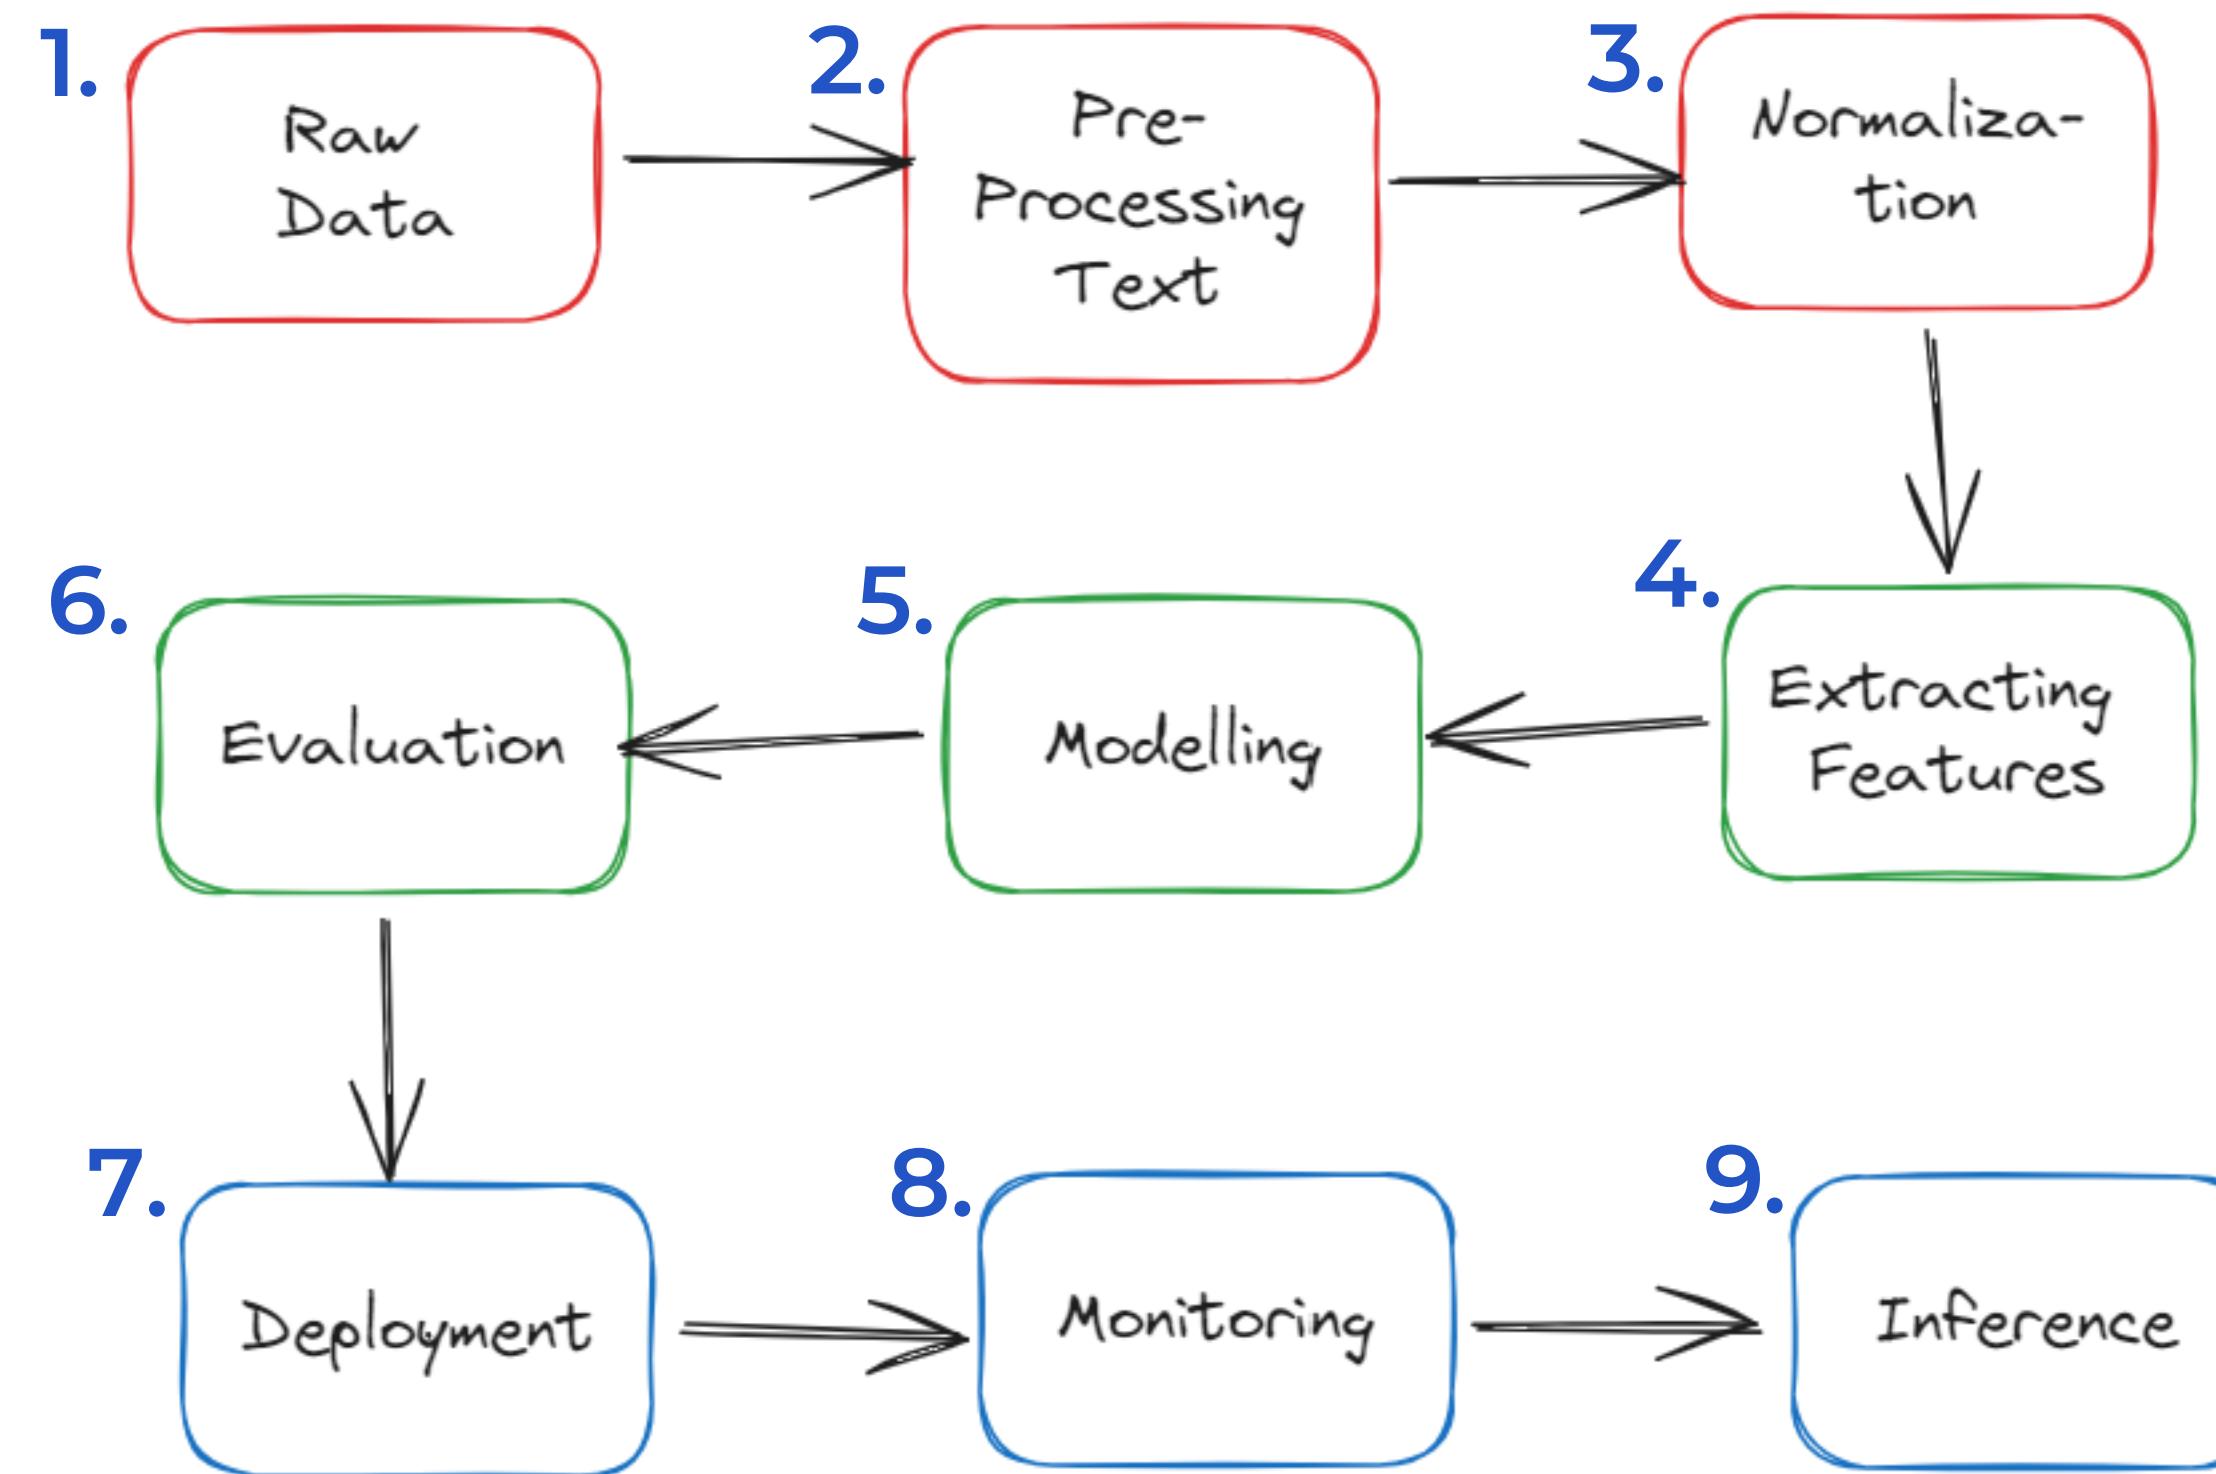
# Content

---

## Lecture 2: Feature Extraction and Word Representations

- Bag of Words, TF-IDF
- One-hot vectors, Count-Based Methods
- Word2Vec, Glove, FastText
- Evaluation and Current State

# NLP pipeline



# Bag of Words

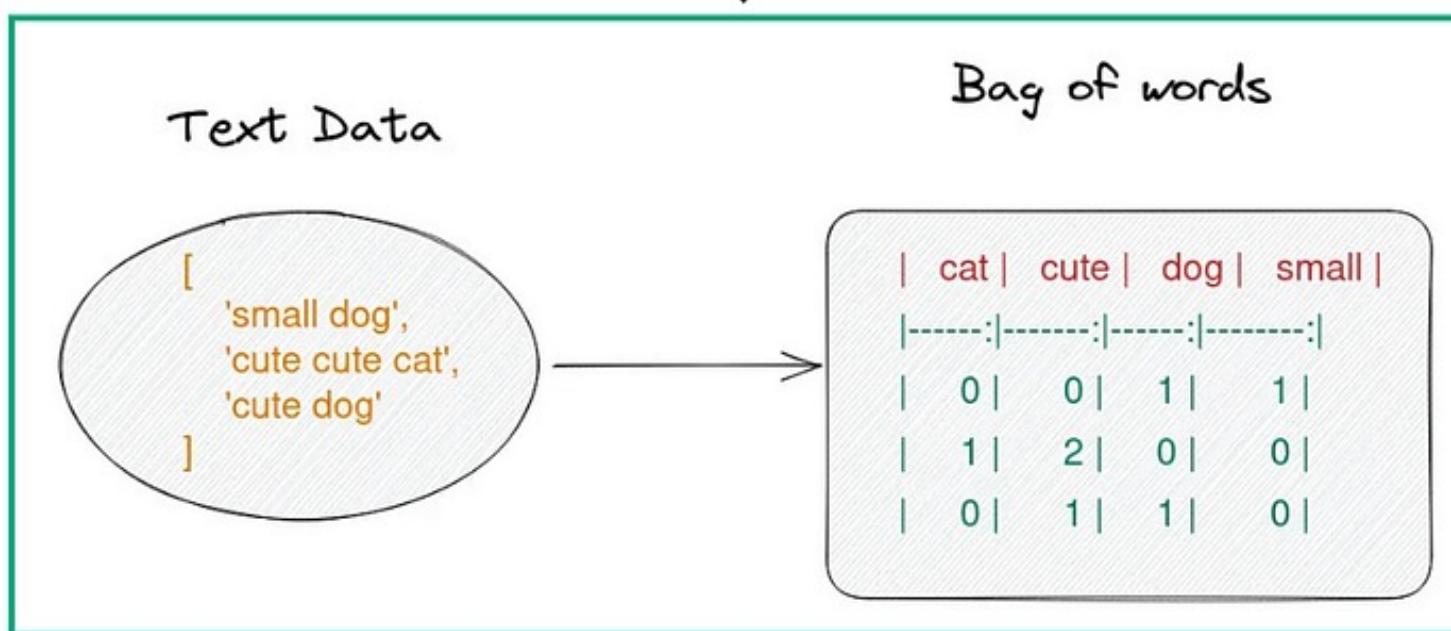
- Feature matrix from classical ML

	Feature-1	Feature-2	Feature-3	Feature-4	...	Feature-n	
Sample-1	$x_1^1$	$x_2^1$	$x_3^1$	$x_4^1$	...	...	$x_n^1$
Sample-2	$x_1^2$	$x_2^2$	$x_3^2$	$x_4^2$	...	...	$x_n^2$
Sample-3	$x_1^3$	$x_2^3$	$x_3^3$	$x_4^3$	...	...	$x_n^3$
...	...	...	...	...	...	...	...
Sample-m	$x_1^m$	$x_2^m$	$x_3^m$	$x_4^m$	...	...	$x_n^m$

# Bag of Words

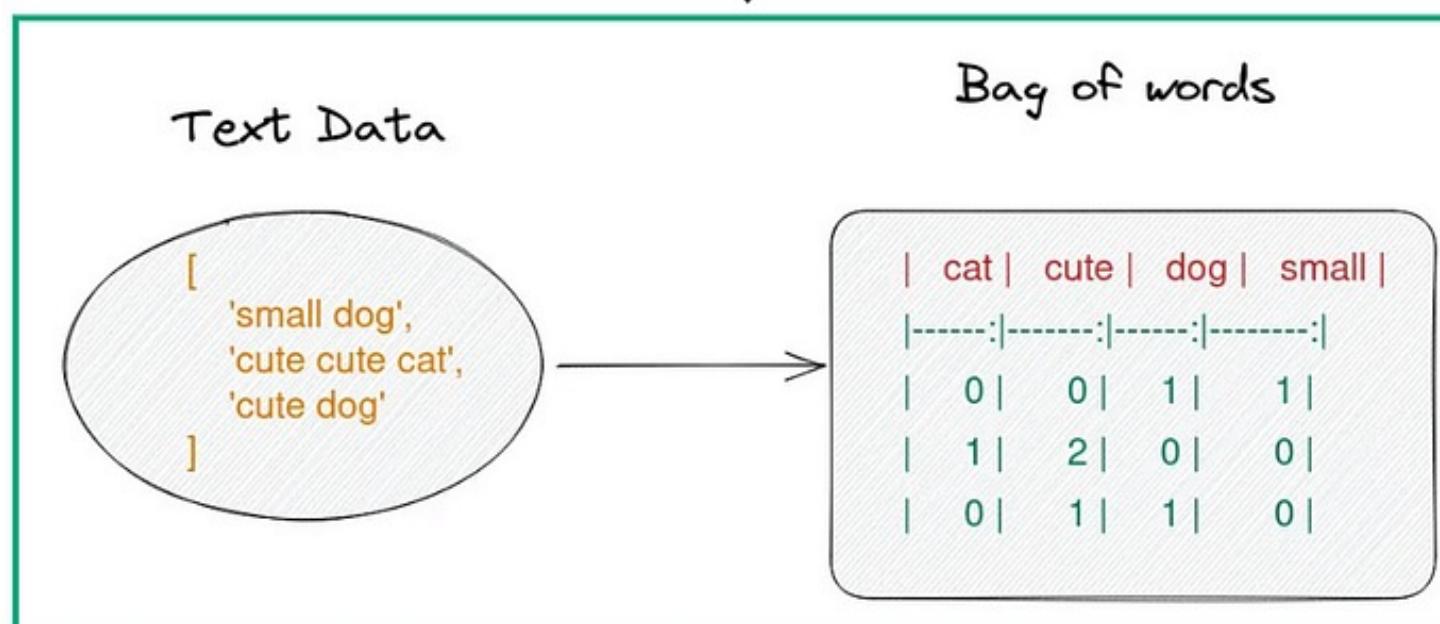
- Feature matrix from classical ML
- A text, such as a paragraph or document, is presented as a set of words, and the frequency with which each word appears in the text is counted

	Feature-1	Feature-2	Feature-3	Feature-4	...	Feature-n	
Sample-1	$x_1^1$	$x_2^1$	$x_3^1$	$x_4^1$	...	...	$x_n^1$
Sample-2	$x_1^2$	$x_2^2$	$x_3^2$	$x_4^2$	...	...	$x_n^2$
Sample-3	$x_1^3$	$x_2^3$	$x_3^3$	$x_4^3$	...	...	$x_n^3$
...	...	...	...	...	...	...	...
Sample-m	$x_1^m$	$x_2^m$	$x_3^m$	$x_4^m$	...	...	$x_n^m$



# Bag of Words

- Feature matrix from classical ML
- A text, such as a paragraph or document, is presented as a set of words, and the frequency with which each word appears in the text is counted



	Feature-1	Feature-2	Feature-3	Feature-4	...	Feature-n	
Sample-1	$x_1^1$	$x_2^1$	$x_3^1$	$x_4^1$	...	...	$x_n^1$
Sample-2	$x_1^2$	$x_2^2$	$x_3^2$	$x_4^2$	...	...	$x_n^2$
Sample-3	$x_1^3$	$x_2^3$	$x_3^3$	$x_4^3$	...	...	$x_n^3$
Sample-m	$x_1^m$	$x_2^m$	$x_3^m$	$x_4^m$	...	...	$x_n^m$

## Problems:

- Loss of word order → Identical vectors  
"The dog bit the man", "The man bit the dog"
- Ignoring semantics → No connection between synonyms  
"I drive a car", "I drive an automobile"
- The curse of dimensionality

# TF-IDF

---

- Idea: highlight words that appear frequently in this text, but rarely in other texts

# TF-IDF

---

- Idea: highlight words that appear frequently in this text, but rarely in other texts
- Term Frequency - Inverse Document Frequency (TF-IDF):

$$\text{tf}(t, d) = \frac{n_t}{\sum_k n_k}$$

# TF-IDF

---

- Idea: highlight words that appear frequently in this text, but rarely in other texts
- Term Frequency - Inverse Document Frequency (TF-IDF):

$$\text{tf}(t, d) = \frac{n_t}{\sum_k n_k}$$

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

# TF-IDF

---

- Idea: highlight words that appear frequently in this text, but rarely in other texts
- Term Frequency - Inverse Document Frequency (TF-IDF):

$$tf(t, d) = \frac{n_t}{\sum_k n_k}$$

$$tf\text{-}idf(t, d, D) = tf(t, d) \times idf(t, D)$$

$$idf(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

where -  $n_t$  is the number of occurrences of word  $t$  in the document  $d$ ,  $N = |D|$

- Usage → `sklearn.feature_extraction.text import TfidfVectorizer`

# Example: TF-IDF

---

*Sentence A:* The car is driven on the road.

*Sentence B:* The truck is driven on the highway.

Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$		
Car	1/7	0	$\log(2/1)=0.3$		
Truck	0	1/7	$\log(2/1)=0.3$		
Is	1/7	1/7	$\log(2/2)=0$		
Driven	1/7	1/7	$\log(2/2)=0$		
On	1/7	1/7	$\log(2/2)=0$		
The	1/7	1/7	$\log(2/2)=0$		
Road	1/7	0	$\log(2/1)=0.3$		
Highway	0	1/7	$\log(2/1)=0.3$		

# Example: TF-IDF

---

*Sentence A:* The car is driven on the road.

*Sentence B:* The truck is driven on the highway.

Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$		
Car	1/7	0	$\log(2/1)=0.3$		
Truck	0	1/7	$\log(2/1)=0.3$		
Is	1/7	1/7	$\log(2/2)=0$		
Driven	1/7	1/7	$\log(2/2)=0$		
On	1/7	1/7	$\log(2/2)=0$		
The	1/7	1/7	$\log(2/2)=0$		
Road	1/7	0	$\log(2/1)=0.3$		
Highway	0	1/7	$\log(2/1)=0.3$		

Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$	0	0
Car	1/7	0	$\log(2/1)=0.3$	0.043	0
Truck	0	1/7	$\log(2/1)=0.3$	0	0.043
Is	1/7	1/7	$\log(2/2)=0$	0	0
Driven	1/7	1/7	$\log(2/2)=0$	0	0
On	1/7	1/7	$\log(2/2)=0$	0	0
The	1/7	1/7	$\log(2/2)=0$	0	0
Road	1/7	0	$\log(2/1)=0.3$	0.043	0
Highway	0	1/7	$\log(2/1)=0.3$	0	0.043

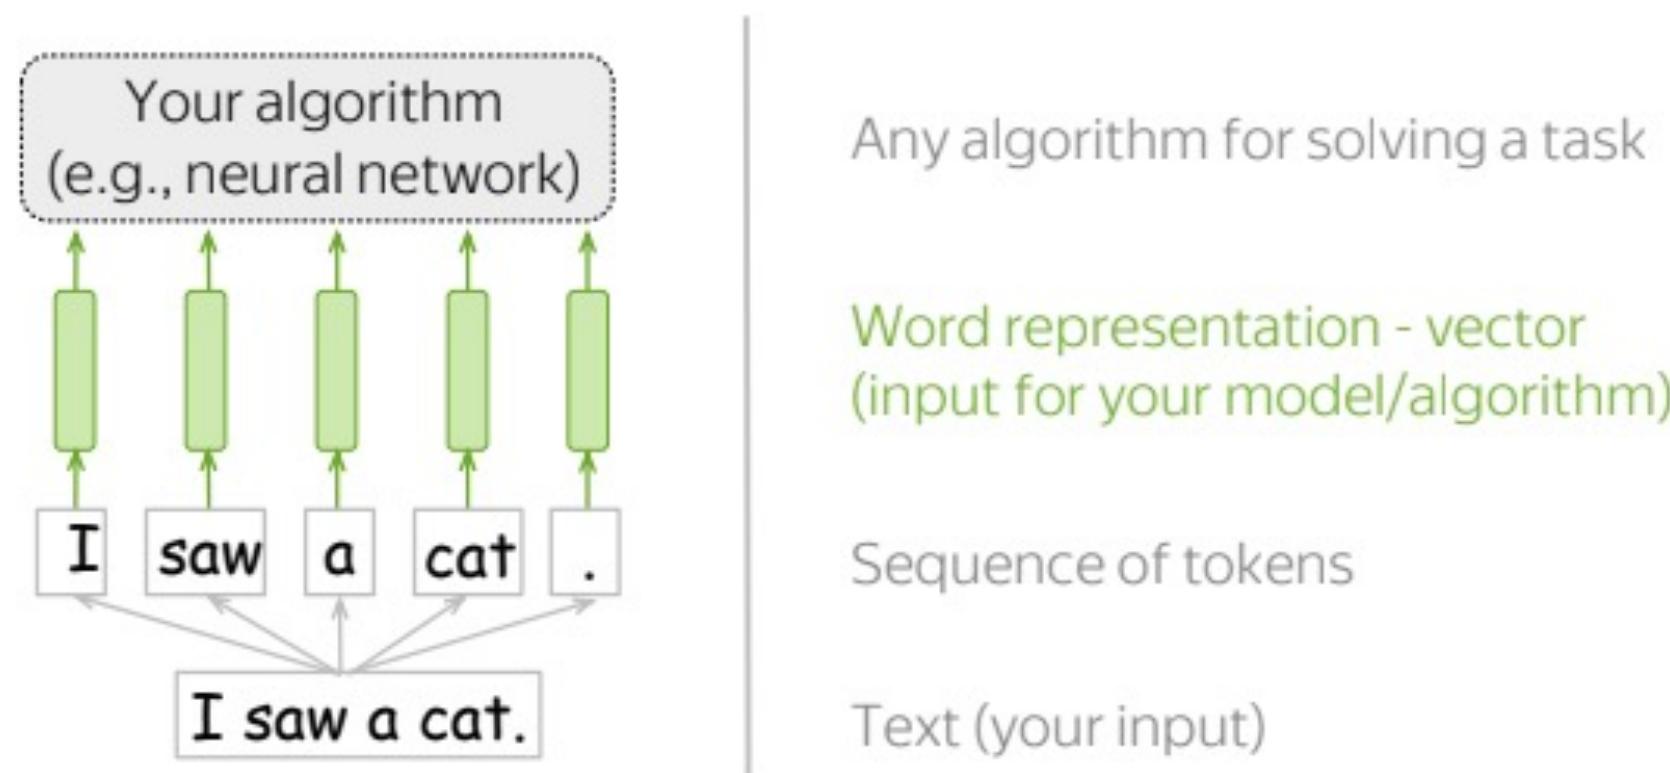
# Aggregation of word representations

---

- Traditional methods treat words as indexes rather than concepts
- Each word  $w \in W$  is associated with a vector  $v_w \in \mathbb{R}^m$  — representation of the word (word embedding),  $m$  — dimension of the space

# Aggregation of word representations

- Traditional methods treat words as indexes rather than concepts
- Each word  $w \in W$  is associated with a vector  $v_w \in \mathbb{R}^m$  — representation of the word (word embedding),  $m$  — dimension of the space
- We will calculate the representation of the document as an aggregate function of the vectors of the document's words



# One-hot Vectors

---

- Each word  $w \in W$  corresponds to a one-hot vector:

$$v_w = [0, \dots, 0, 1, 0, \dots, 0] \in \mathbb{R}^{|W|}$$

# One-hot Vectors

---

- Each word  $w \in W$  corresponds to a one-hot vector:

$$v_w = [0, \dots, 0, 1, 0, \dots, 0] \in \mathbb{R}^{|W|}$$

One is 1, the rest are 0



dog      0...0...0**1**0....0...0

cat      0...0**1**0...0....0...0

table    0...0...0....00**1**0...



Embedding dimension =  
vocabulary size

## Problems:

- Sparseness
- Large dimensionality
- Orthogonality of all word representations
- No mechanism for processing unfamiliar words (out of vocabulary, OOV) on the test
- Vectors know nothing about meaning  
“cat is as close to dog as it is to table!”

# What is meaning?

---

- Do you know what the word **tezgüino** means ?

A bottle of **tezgüino** is on the table.

Everyone likes **tezgüino**.

**Tezgüino** makes you drunk.

We make **tezgüino** out of corn.

# What is meaning?

---

- Do you know what the word **tezgüino** means ?

A bottle of **tezgüino** is on the table.

Everyone likes **tezgüino**.

**Tezgüino** makes you drunk.

We make **tezgüino** out of corn.



**Tezgüino** is a kind of alcoholic beverage made from corn.

# What is meaning?

---

- Do you know what the word **tezgüino** means ?

A bottle of **tezgüino** is on the table.

Everyone likes **tezgüino**.

**Tezgüino** makes you drunk.

We make **tezgüino** out of corn.



**Tezgüino** is a kind of alcoholic beverage made from corn.

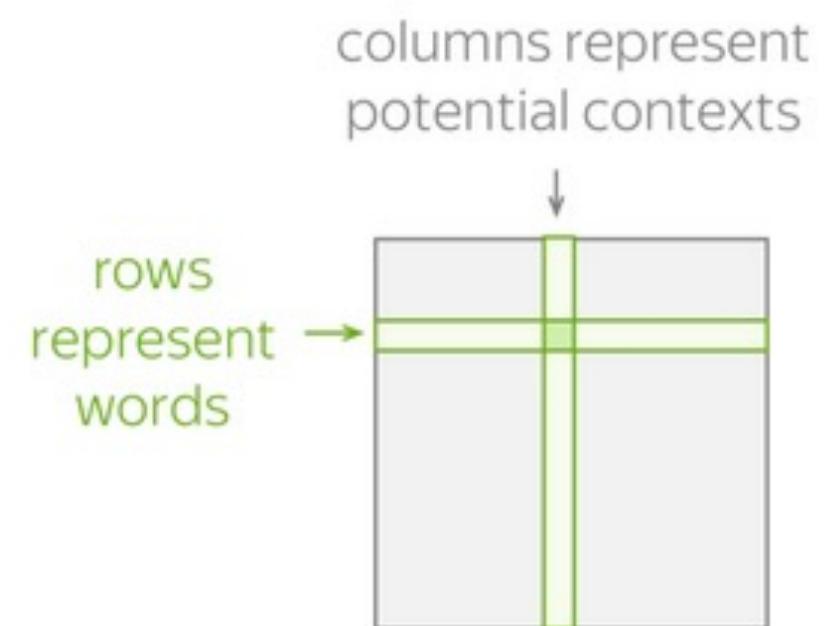
Distributional Hypothesis (Harris 1954, Firth 1957)

Words which frequently appear in similar contexts have similar meaning.

# Count-based methods

---

- We have to put information about contexts into word vectors

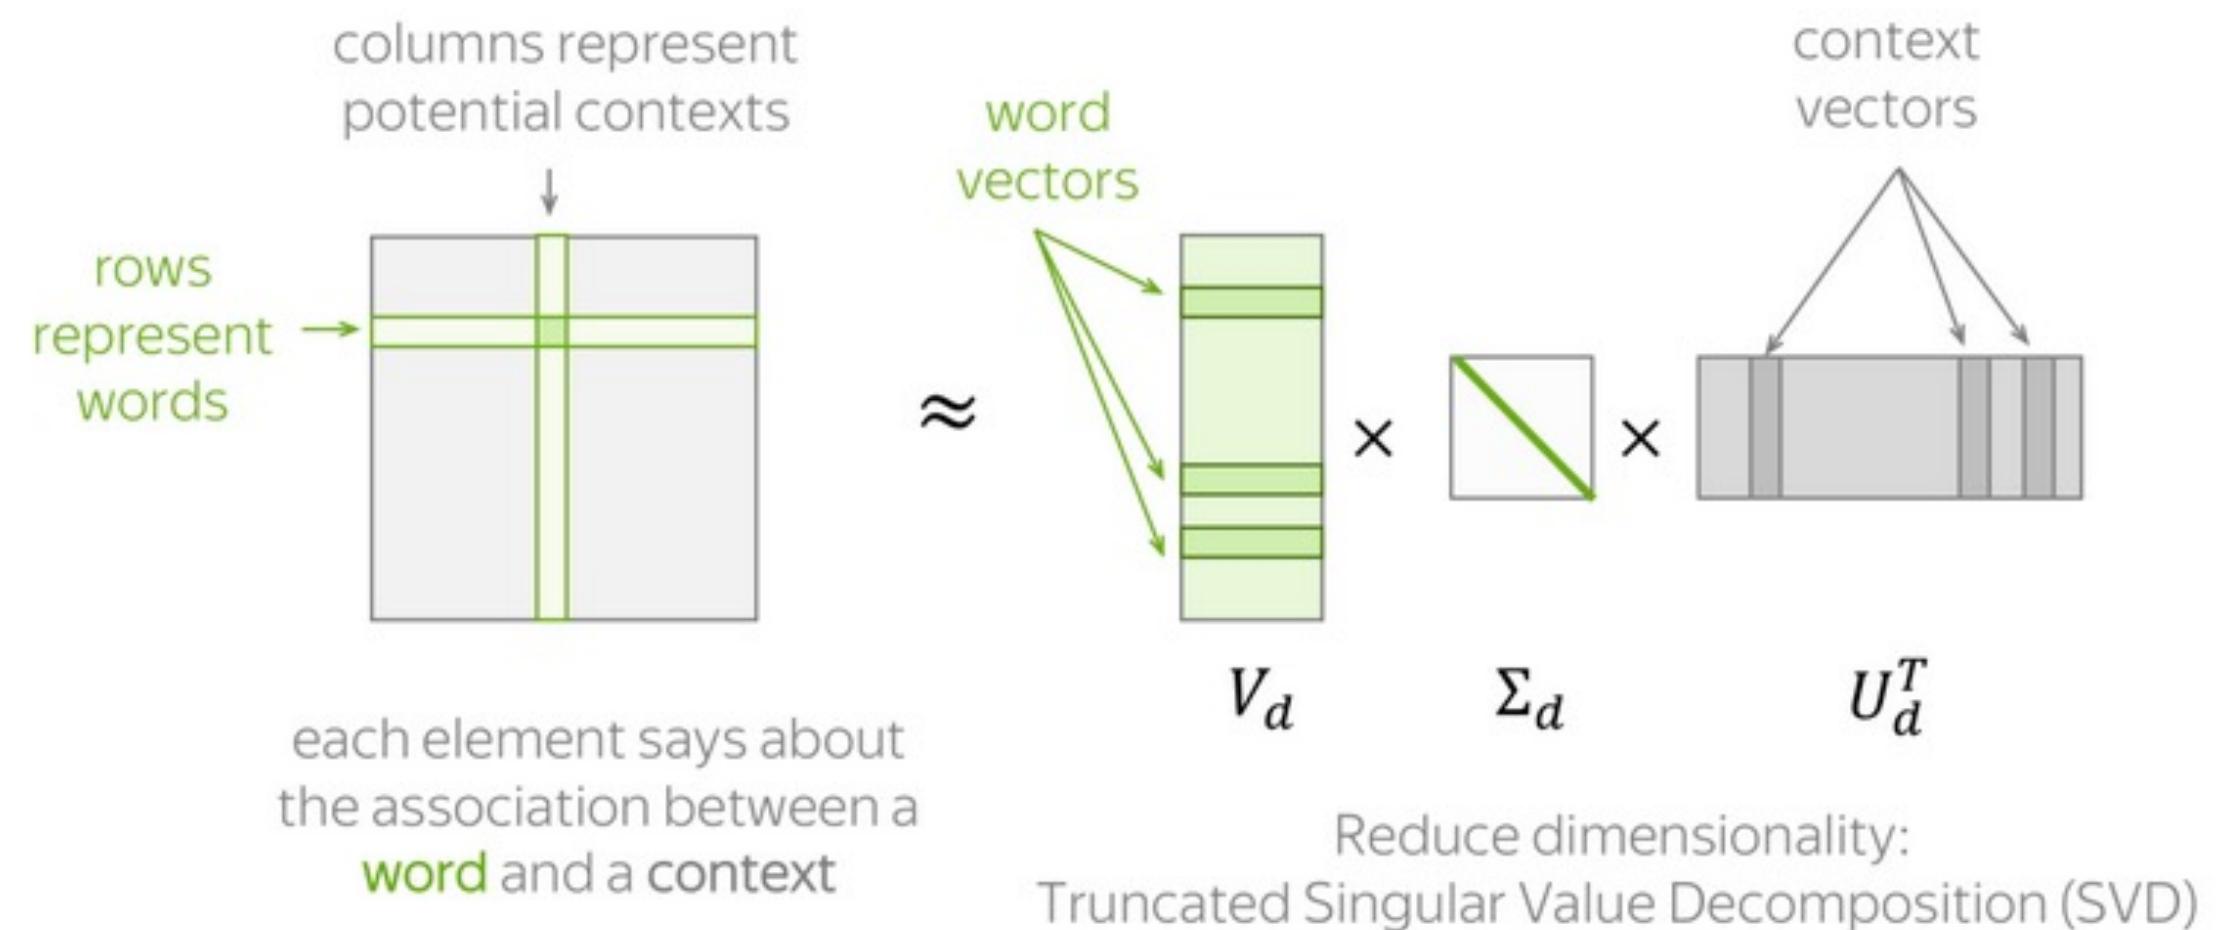


each element says about  
the association between a  
**word** and a context

- How to obtain  $v_w \in \mathbb{R}^m$ ?

# SVD for representations

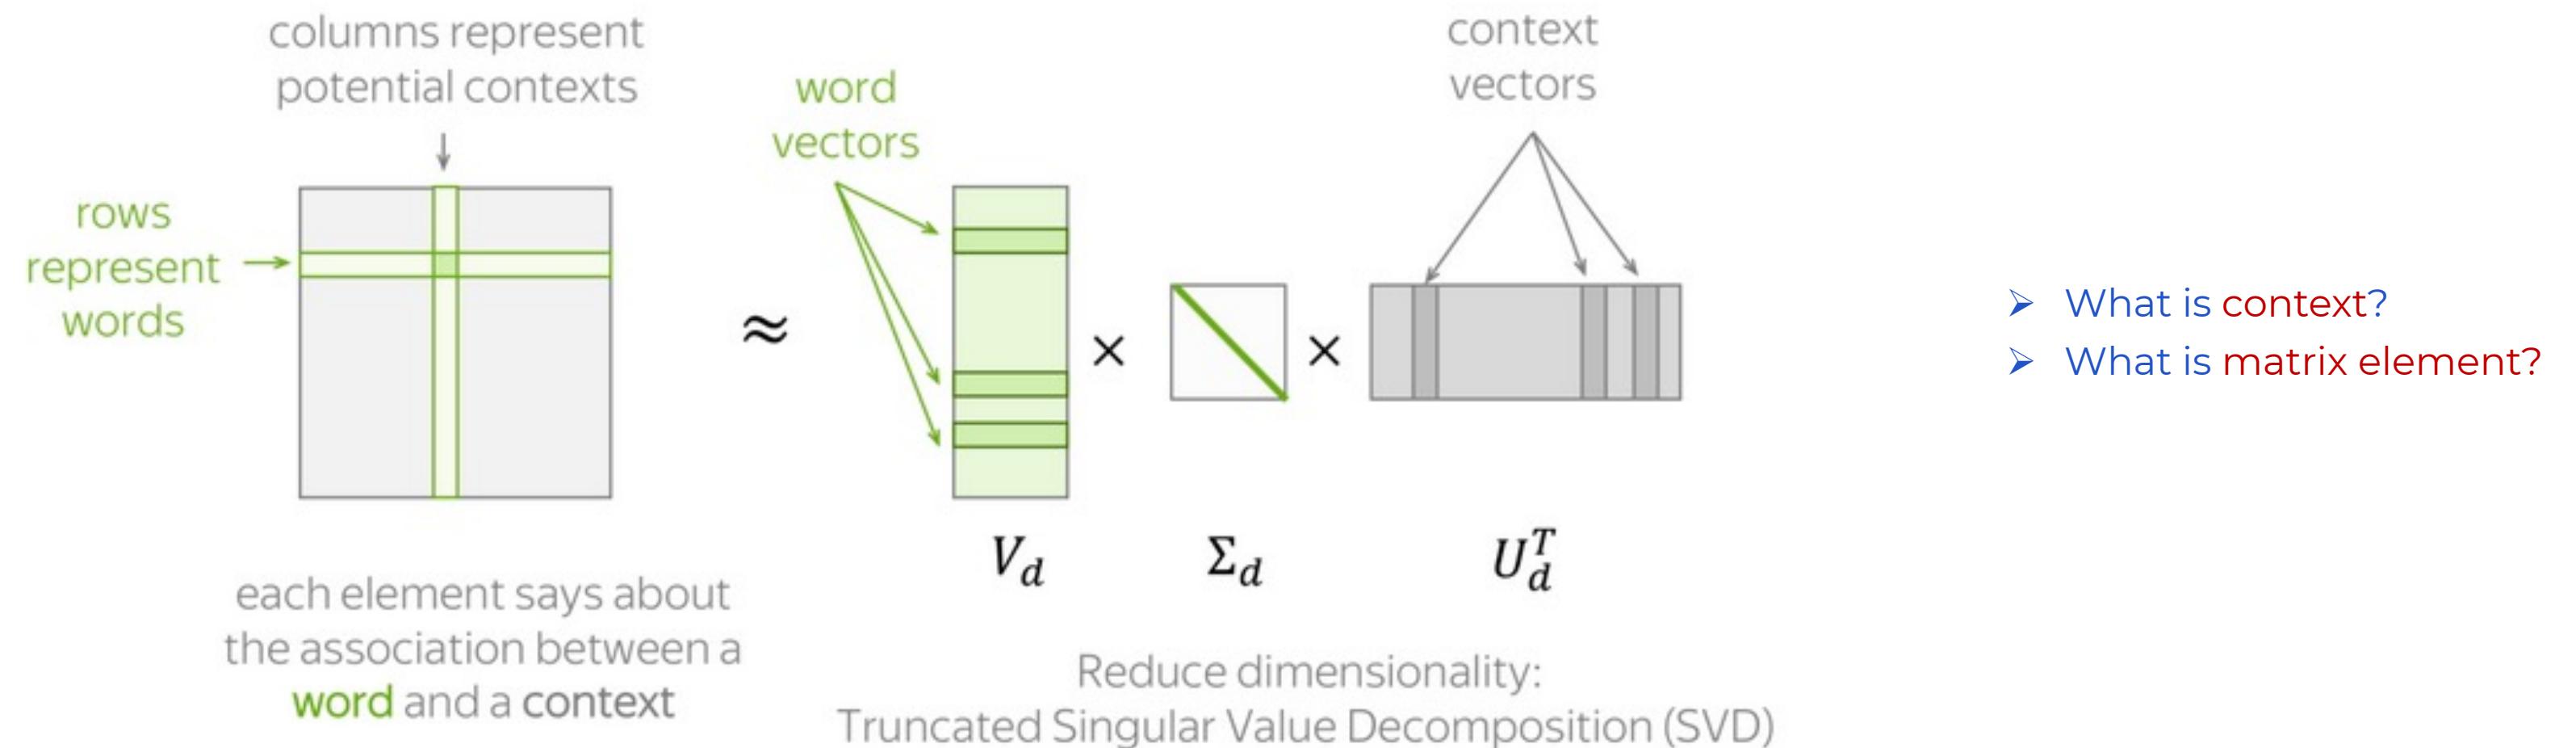
- Word representations via matrix factorization



Initial matrix for the collection  $X \in \mathbb{R}^{|W| \times |W|}$ , approximation  $X \approx \underbrace{V_d}_{\text{word vectors}} \cdot \underbrace{\Sigma_d}_{\text{diagonal}} \cdot \underbrace{U_d^T}_{\text{context vectors}} \rightarrow V_d \approx X U_d \Sigma_d^{-1}$

# SVD for representations

- Word representations via matrix factorization

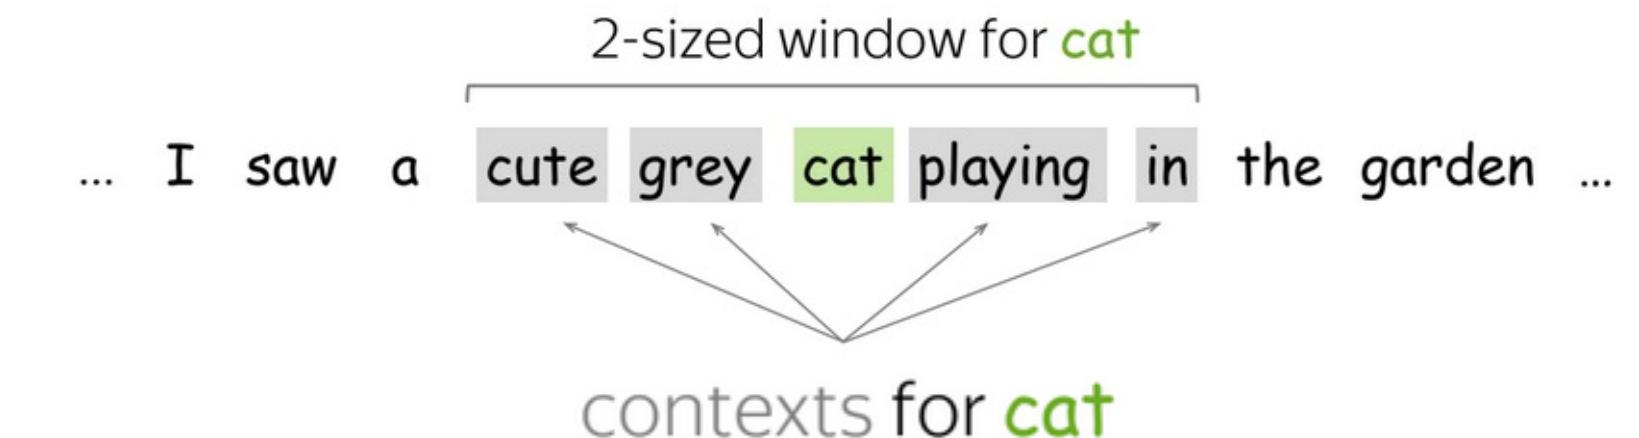


Initial matrix for the collection  $X \in \mathbb{R}^{|W| \times |W|}$ , approximation  $X \approx \underbrace{V_d}_{\text{word vectors}} \cdot \underbrace{\Sigma_d}_{\text{diagonal}} \cdot \underbrace{U_d^T}_{\text{context vectors}}$   $\rightarrow V_d \approx X U_d \Sigma_d^{-1}$

# Collocations

---

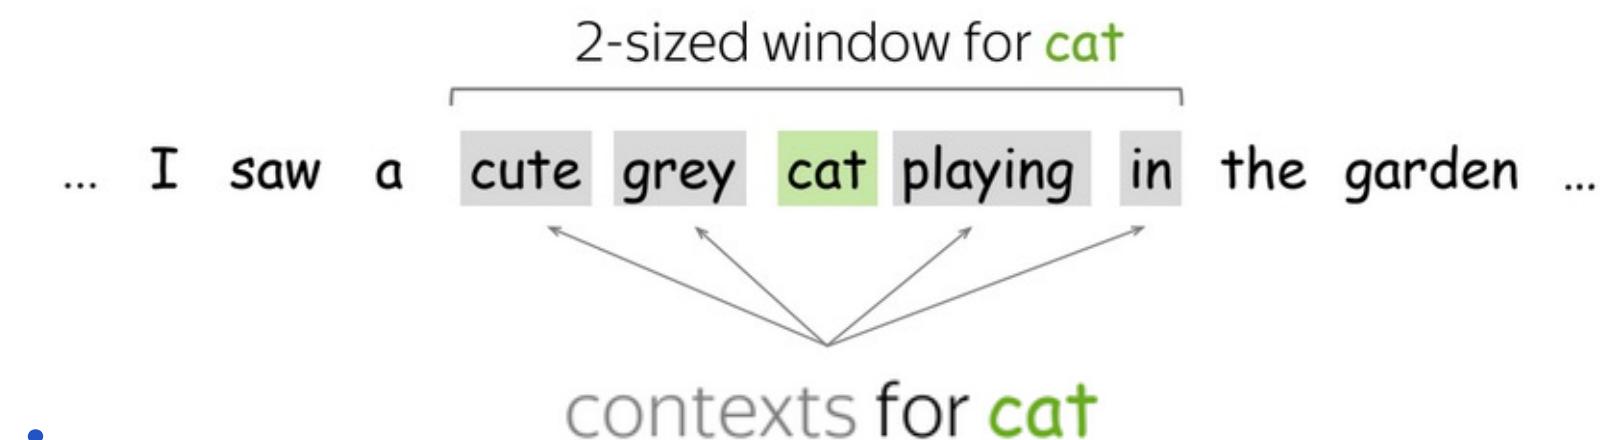
- Co-Occurrence Counts: **context** – surrounding window, **matrix element** - number of times word appears in context



# Collocations

---

- Co-Occurrence Counts: **context** – surrounding window, **matrix element** - number of times word appears in context



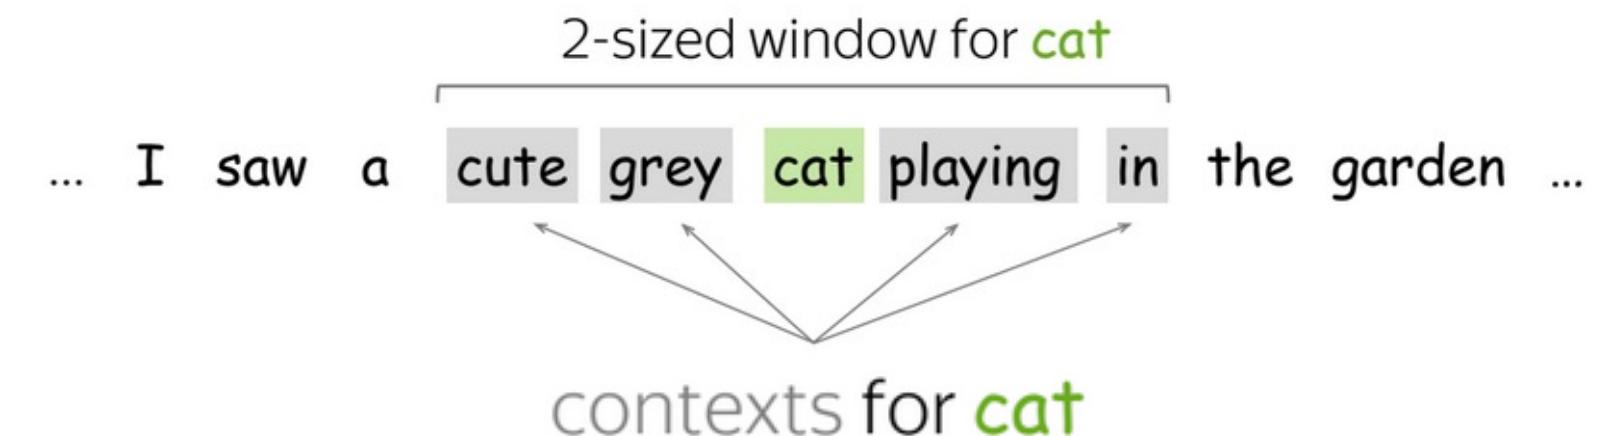
- Positive Pointwise Mutual Information (PPMI): **context** – surrounding window, **matrix element**:

$$\text{PPMI}(w, c) = \max(0, \text{PMI}(w, c)), \quad \text{PMI}(w, c) = \log \frac{P(w, c)}{P(w)P(c)} = \log \frac{N(w, c)|w, c|}{N(w)N(c)}$$

# Collocations

---

- Co-Occurrence Counts: **context** – surrounding window, **matrix element** - number of times word appears in context



- Positive Pointwise Mutual Information (PPMI): **context** – surrounding window, **matrix element**:

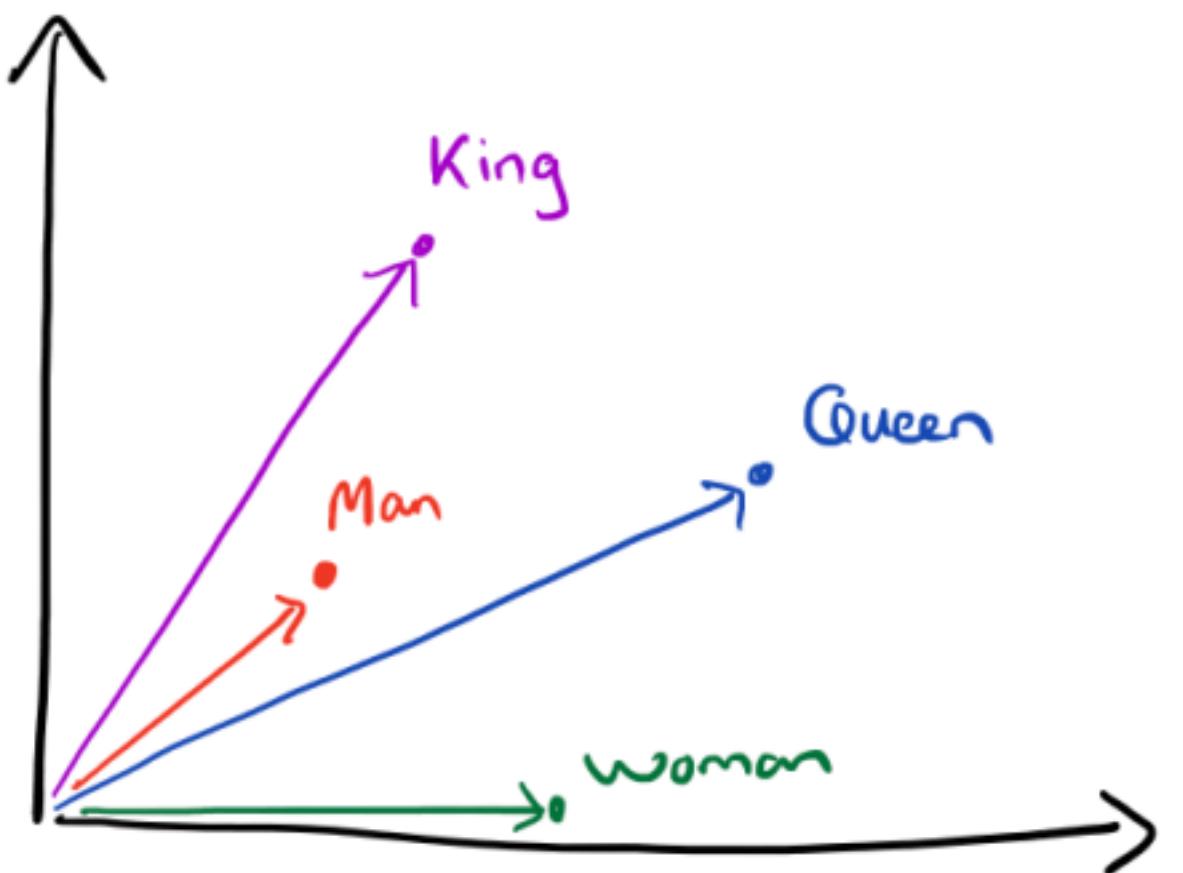
$$\text{PPMI}(w, c) = \max(0, \text{PMI}(w, c)), \quad \text{PMI}(w, c) = \log \frac{P(w, c)}{P(w)P(c)} = \log \frac{N(w, c)|N(w, c)|}{N(w)N(c)}$$

- Latent Semantic Analysis (LSA): **context** – document from collection, **matrix element** – TF-IDF(**w**, **d**, **D**)

# Prediction-based methods

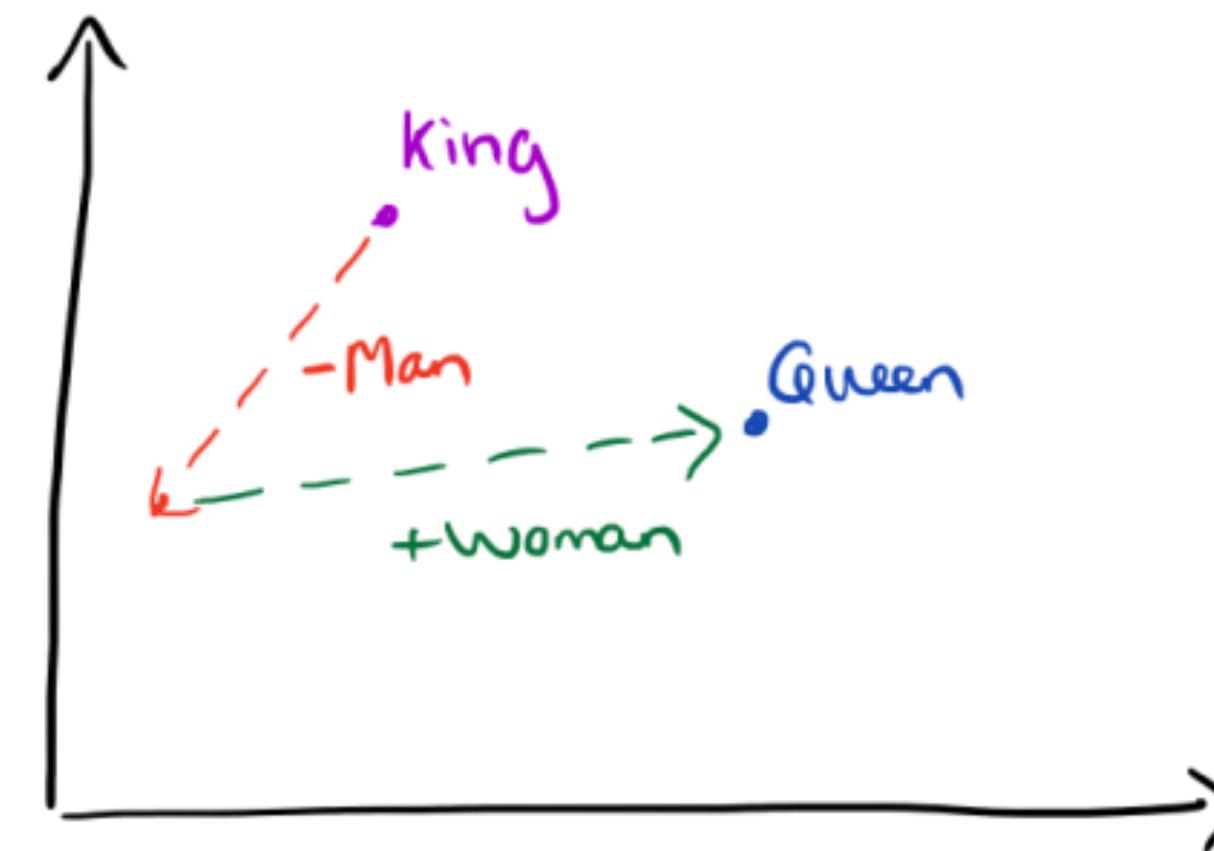
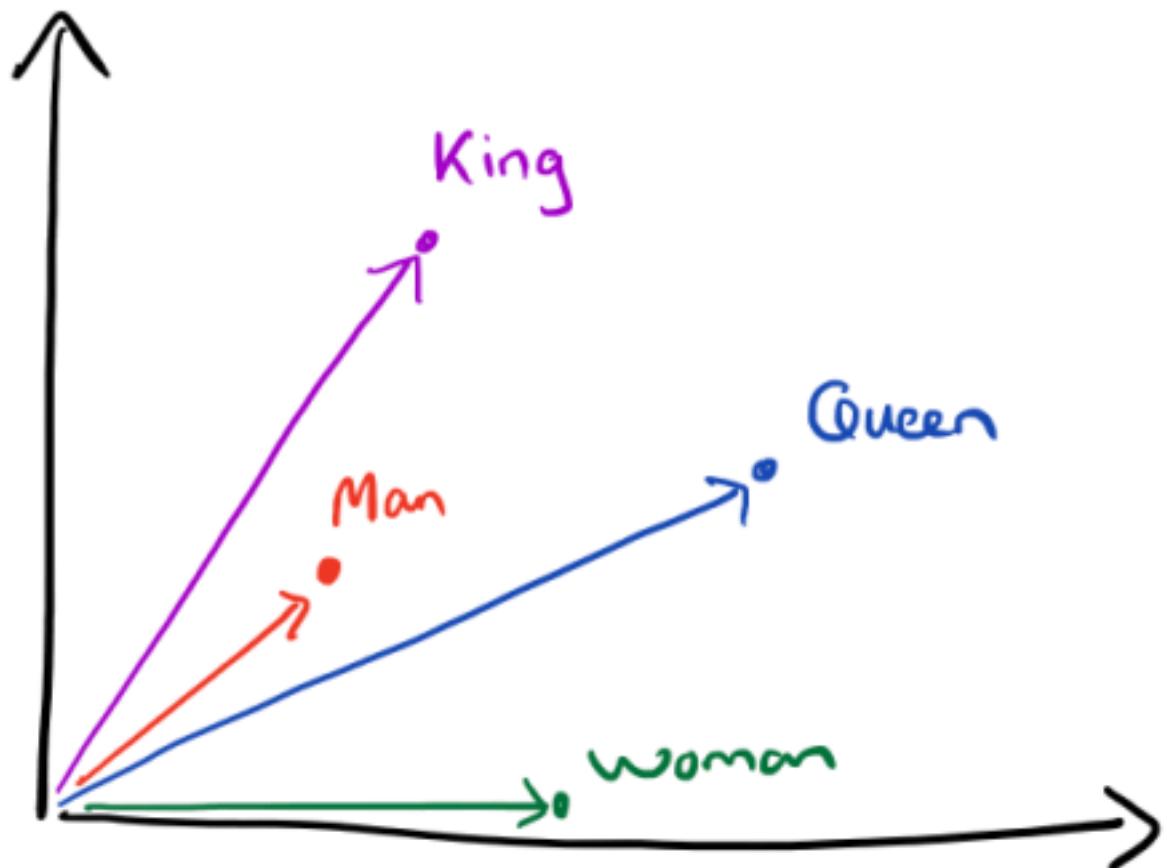
---

- It is possible to learn word vectors that are able to capture the relationships between words



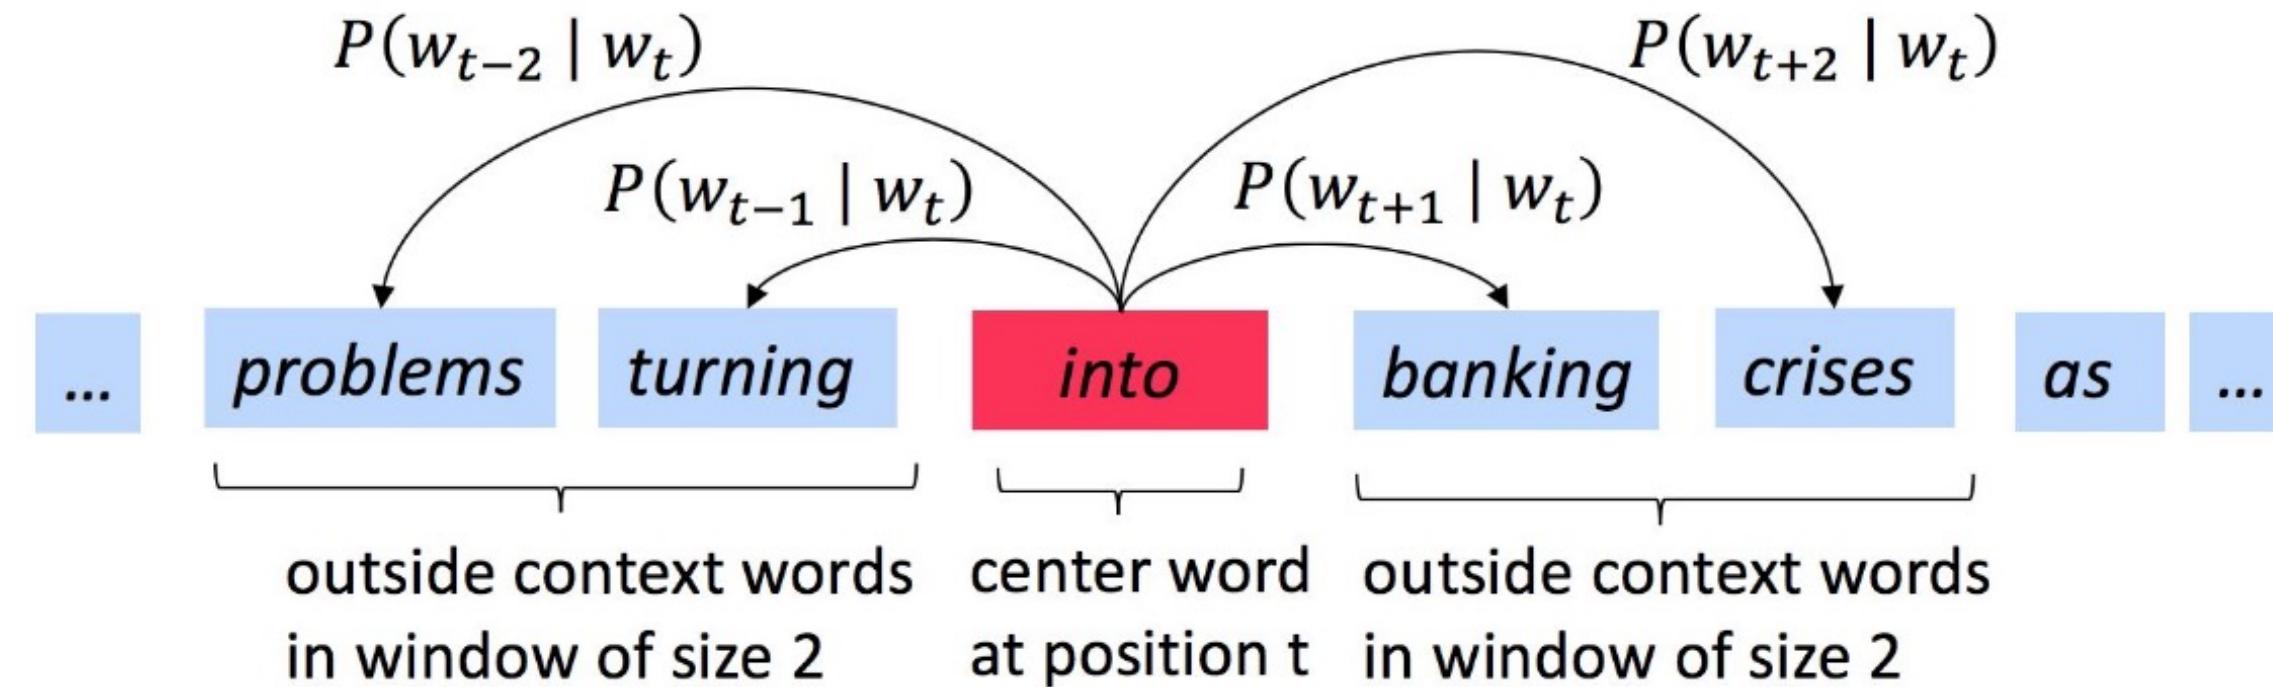
# Prediction-based methods

- It is possible to learn word vectors that are able to capture the relationships between words

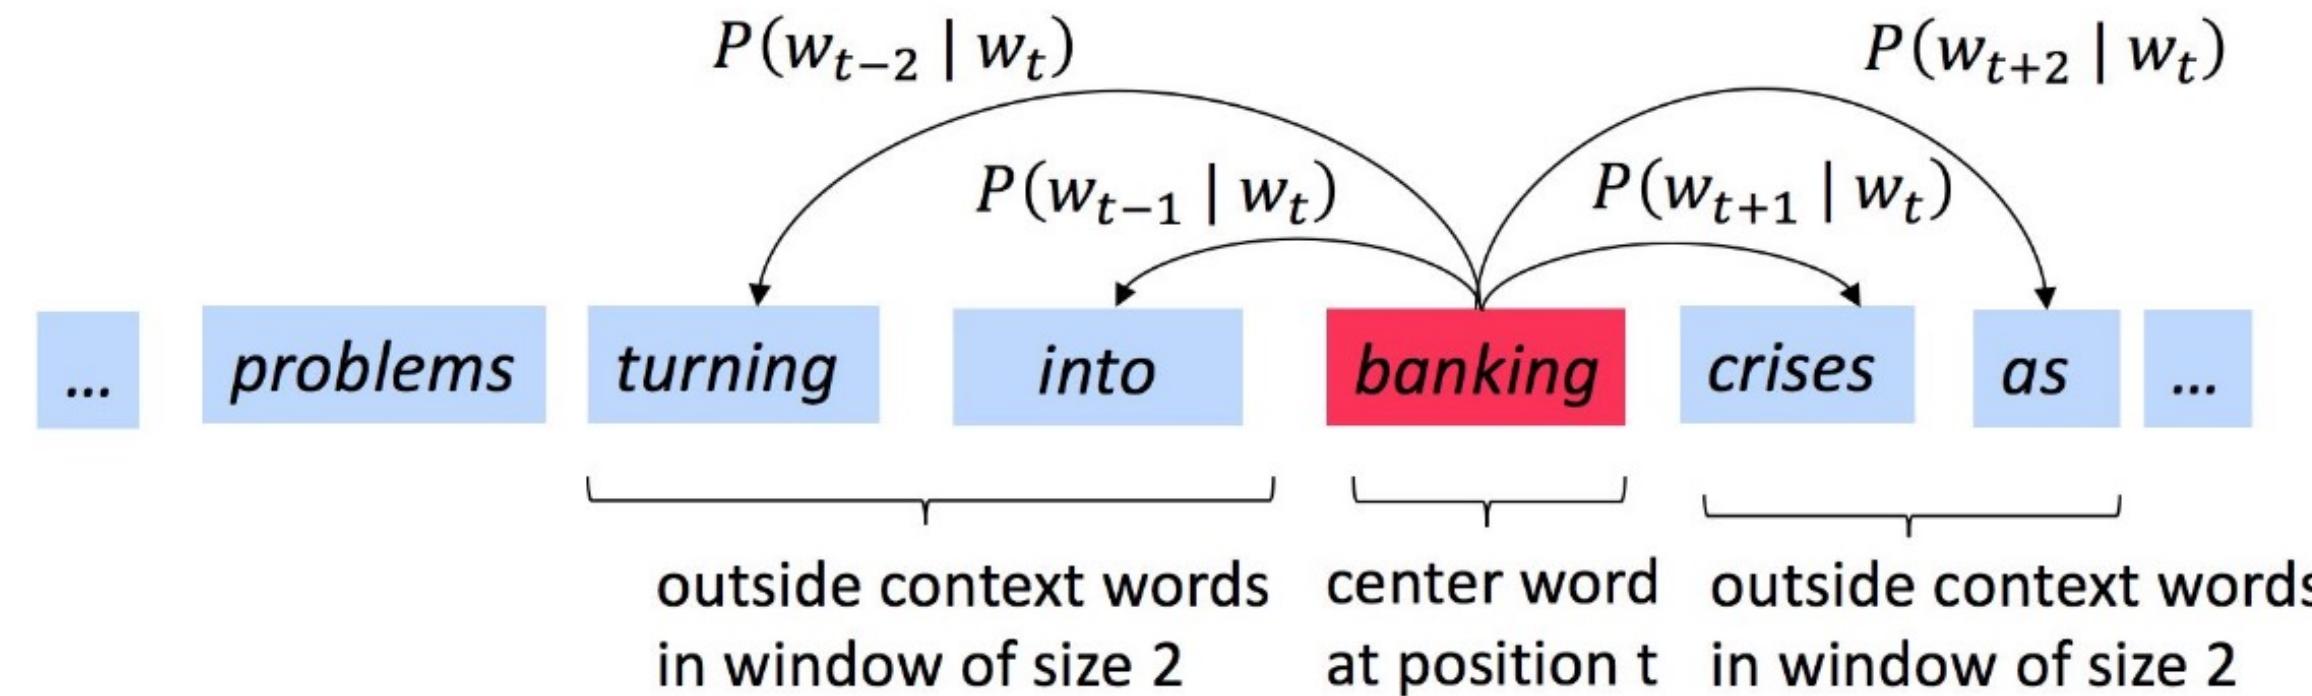
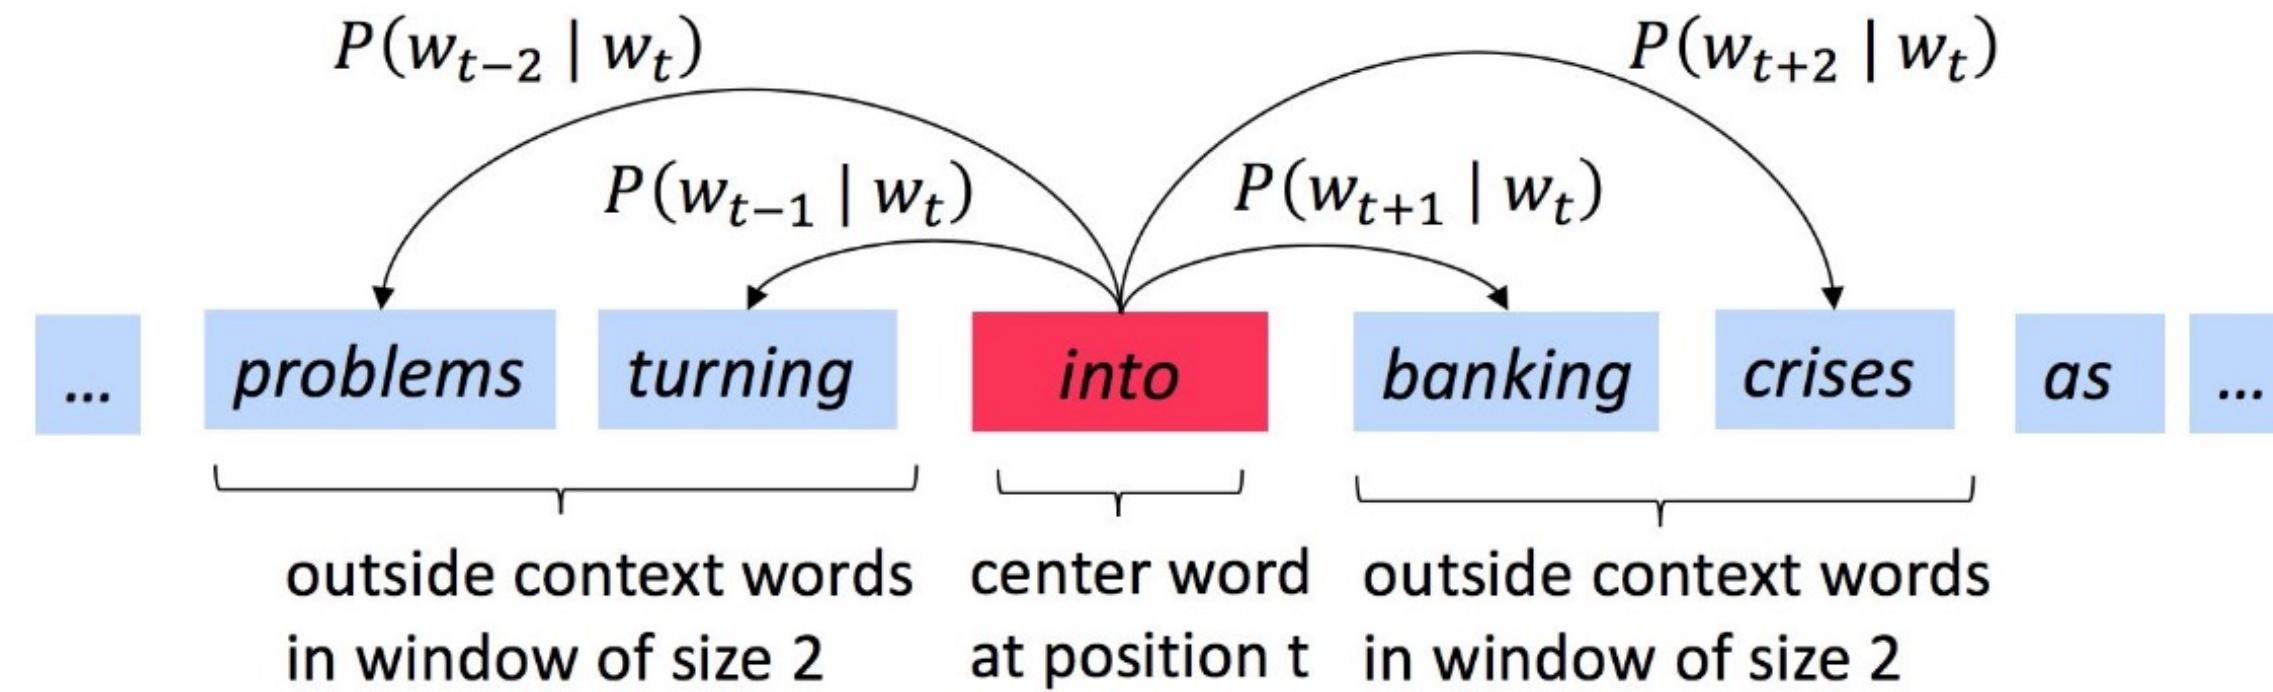


- Word2vec (Mikolov et al. 2013) - a framework for learning word embeddings

# Word2Vec: Idea



# Word2Vec: Idea



# Word2Vec: Objective

---

> Maximize the data likelihood:

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m, \\ j \neq 0}} P(w_{t+j} | \mathbf{w}_t, \theta)$$

# Word2Vec: Objective

---

> Maximize the data likelihood:

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m, \\ j \neq 0}} P(w_{t+j} | w_t, \theta)$$

> Uses negative (log-)likelihood as its loss function:

$$\text{Loss} = J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m, \\ j \neq 0}} \log P(w_{t+j} | w_t, \theta)$$

# Word2Vec: Objective

> Maximize the data likelihood:

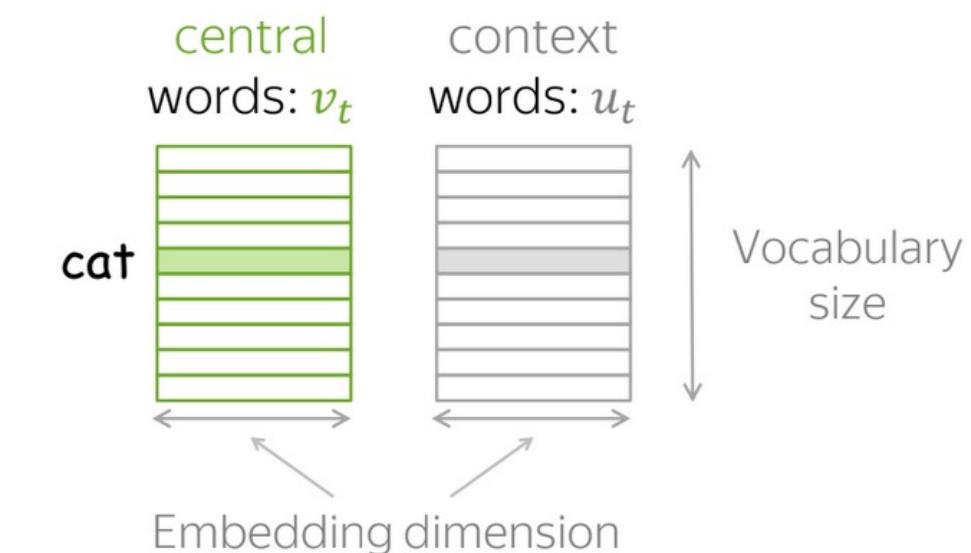
$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m, \\ j \neq 0}} P(w_{t+j} | \mathbf{w}_t, \theta)$$

> Uses negative (log-)likelihood as its loss function:

$$\text{Loss} = J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m, \\ j \neq 0}} \log P(w_{t+j} | \mathbf{w}_t, \theta)$$

> For each word  $w$ , we will have two vectors:  $v_w$  when it is a central word,  $u_w$  when it is a context word:

$$P(o | c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$



# Word2Vec: Objective

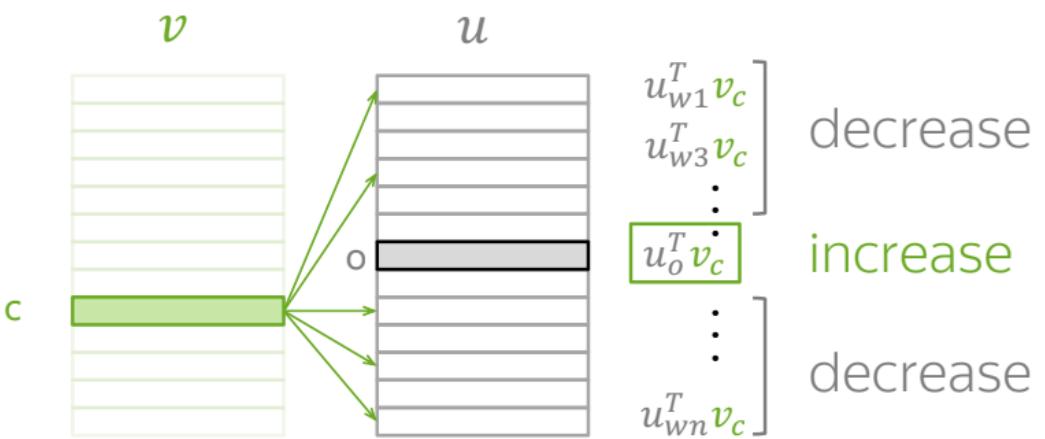
> Maximize the data likelihood:

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m, \\ j \neq 0}} P(w_{t+j} | \mathbf{w}_t, \theta)$$

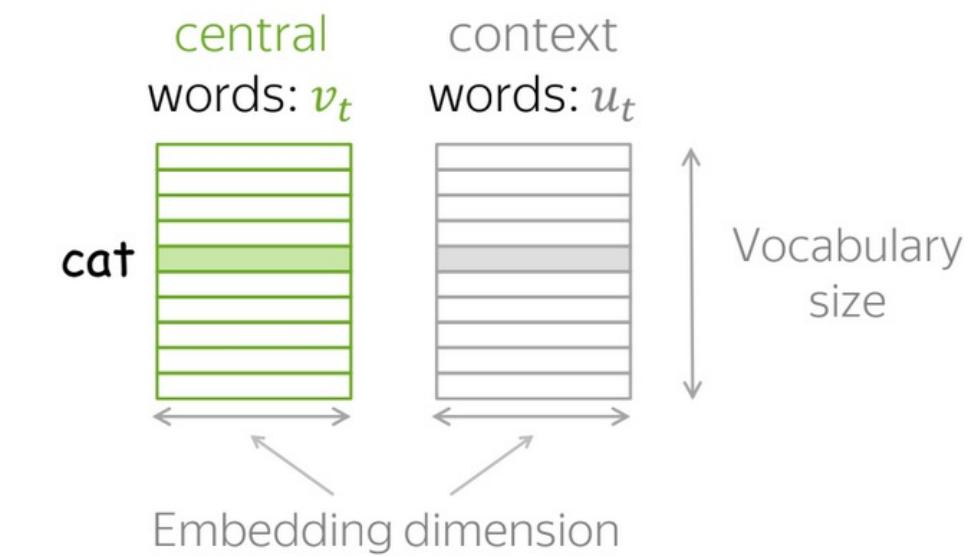
> Uses negative (log-)likelihood as its loss function:

$$\text{Loss} = J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m, \\ j \neq 0}} \log P(w_{t+j} | \mathbf{w}_t, \theta)$$

> For each word  $w$ , we will have two vectors:  $v_w$  when it is a central word,  $u_w$  when it is a context word:



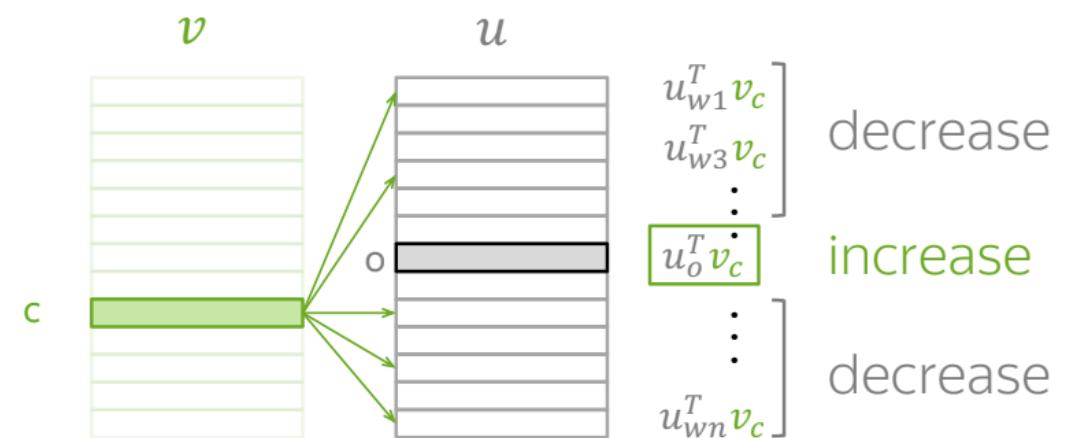
$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$



# Word2Vec: Improvements

---

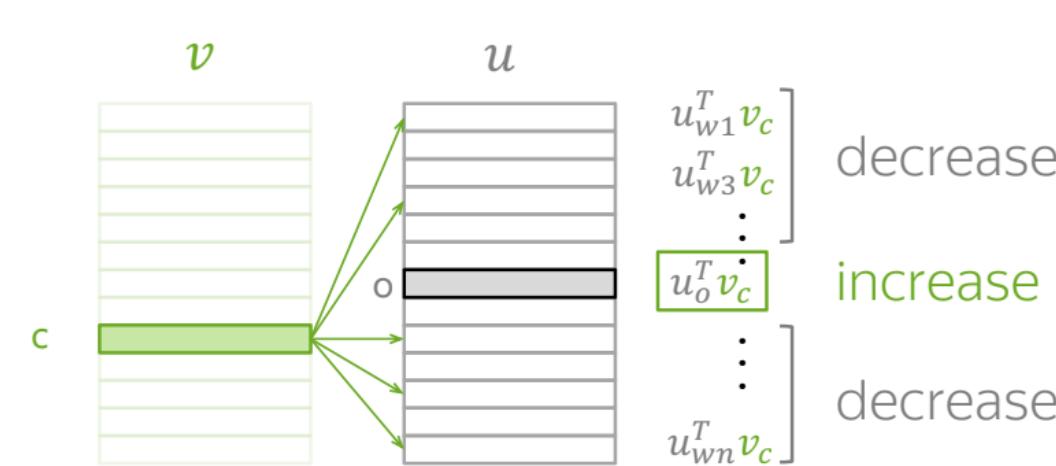
## > Negative sampling:



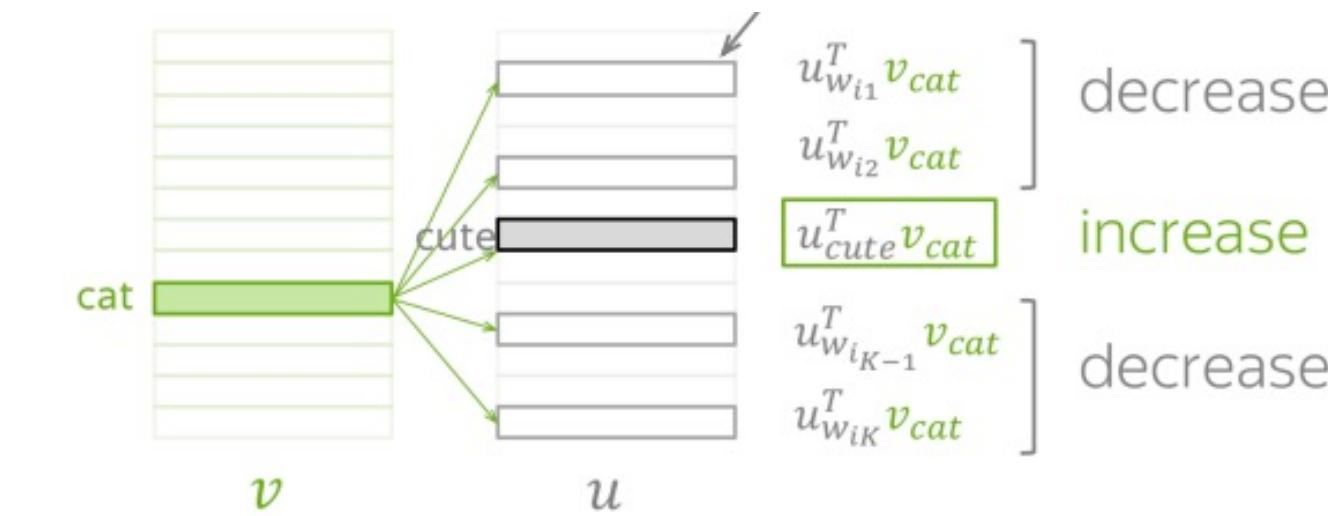
Need to update  $|W| + 1$  vector

# Word2Vec: Improvements

## > Negative sampling:



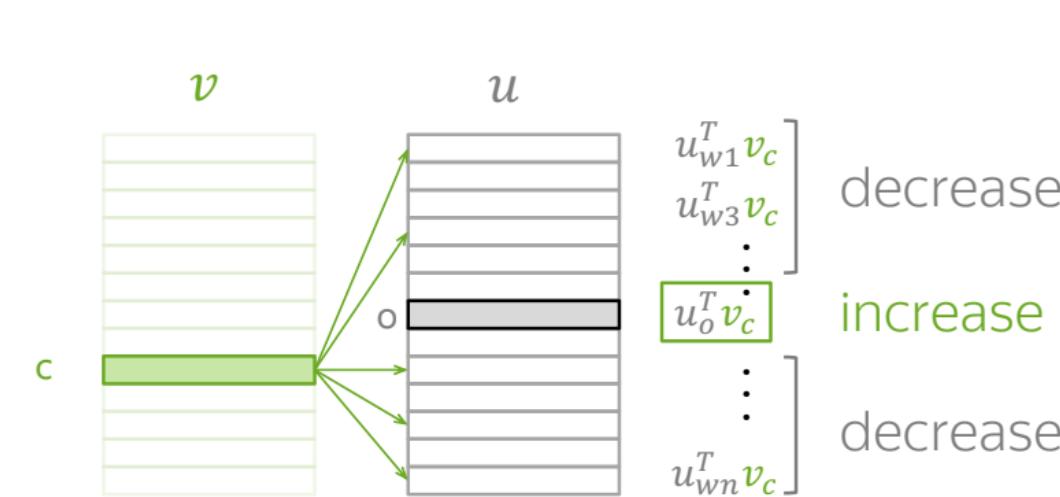
Need to update  $|W| + 1$  vector



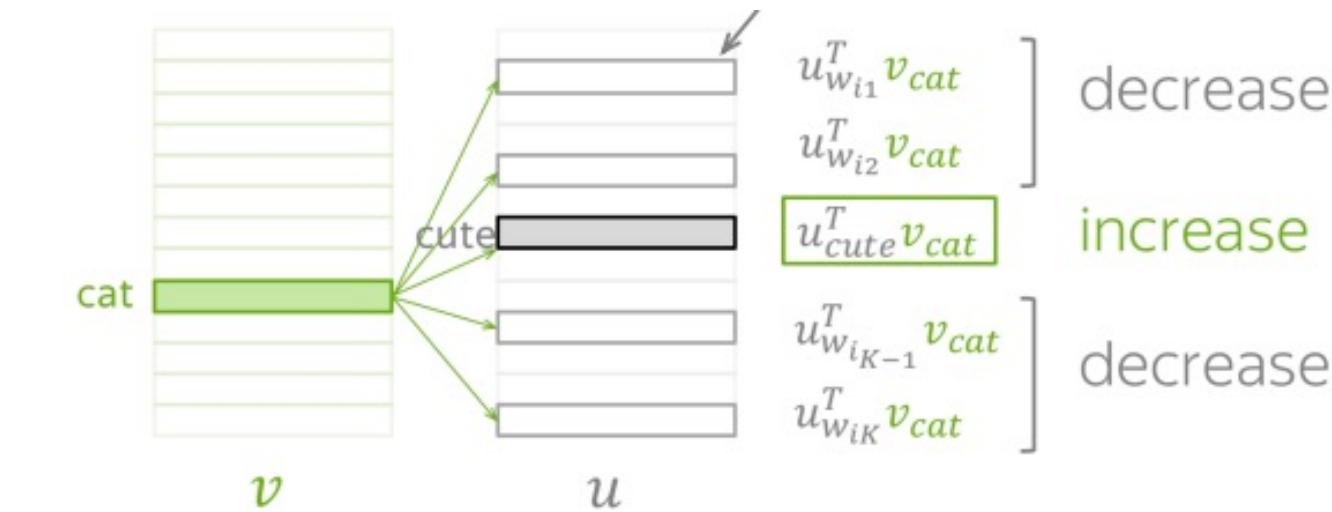
Randomly select K words

# Word2Vec: Improvements

## > Negative sampling:



Need to update  $|W| + 1$  vector



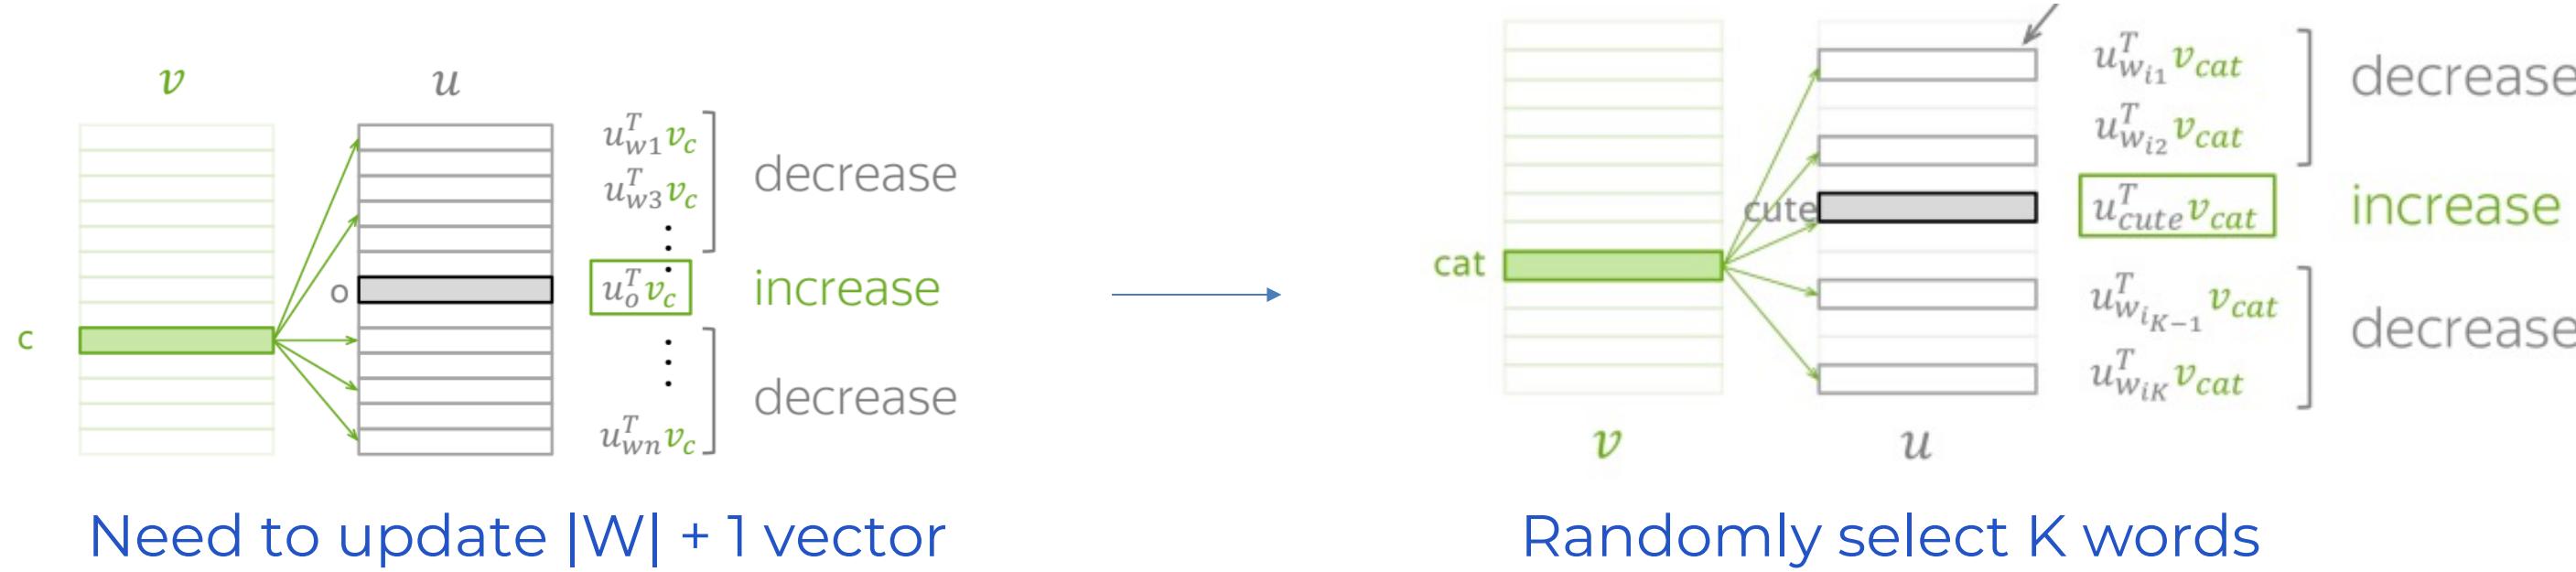
Randomly select  $K$  words

## > Hierarchical softmax:

Idea: replace softmax with another function whose optimisation will have a complexity of  $O(\log |W|)$ .

# Word2Vec: Improvements

## > Negative sampling:



## > Hierarchical softmax:

Idea: replace softmax with another function whose optimisation will have a complexity of  $O(\log |W|)$ .

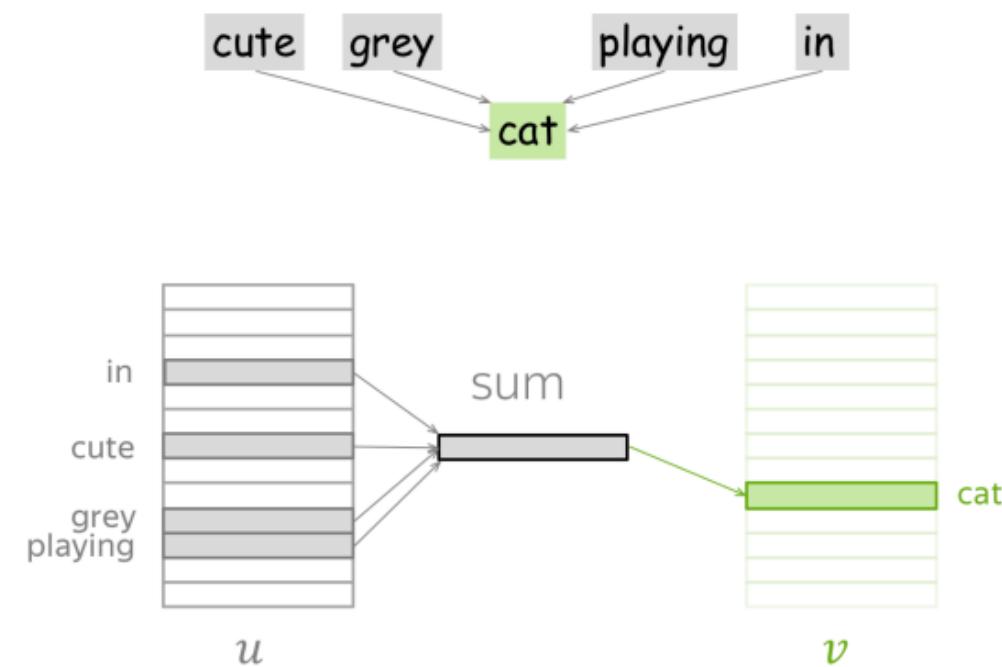
Before training the model on a set of word pairs and their frequencies, a Huffman binary tree is constructed.

Each node of the tree corresponds to a trainable representation. The leaves of the tree correspond to words. The representations in the leaves are the desired representations for the words.

# Word2Vec: Two methods

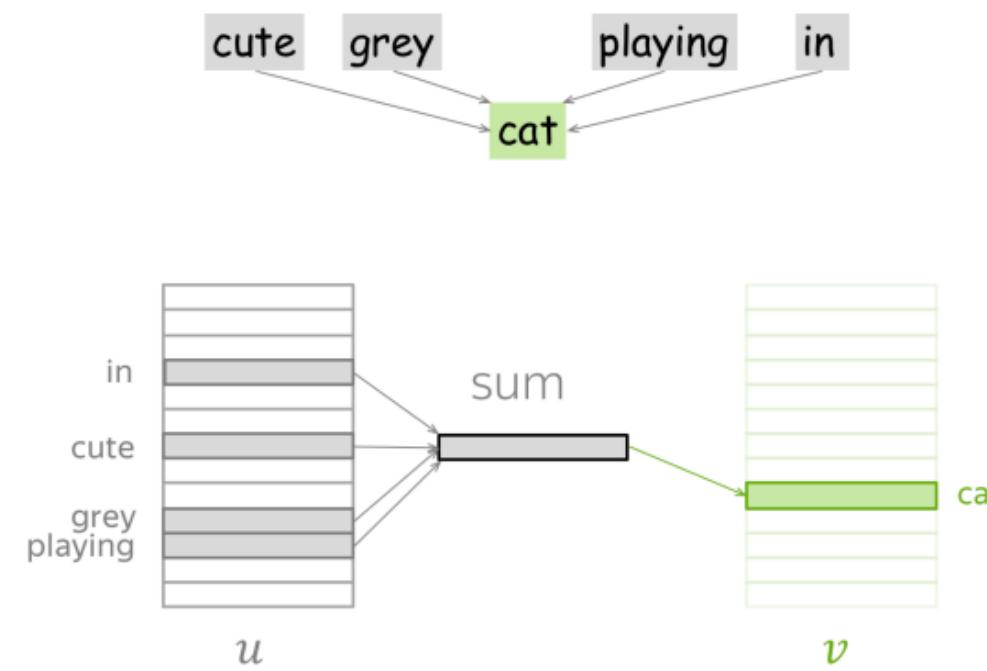
---

## Continuous BOW (CBOW)

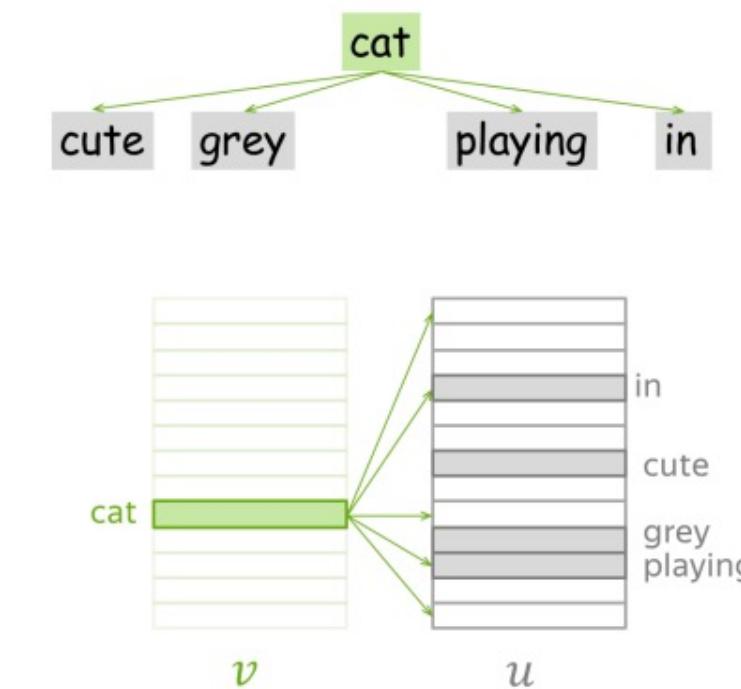


# Word2Vec: Two methods

## Continuous BOW (CBOW)

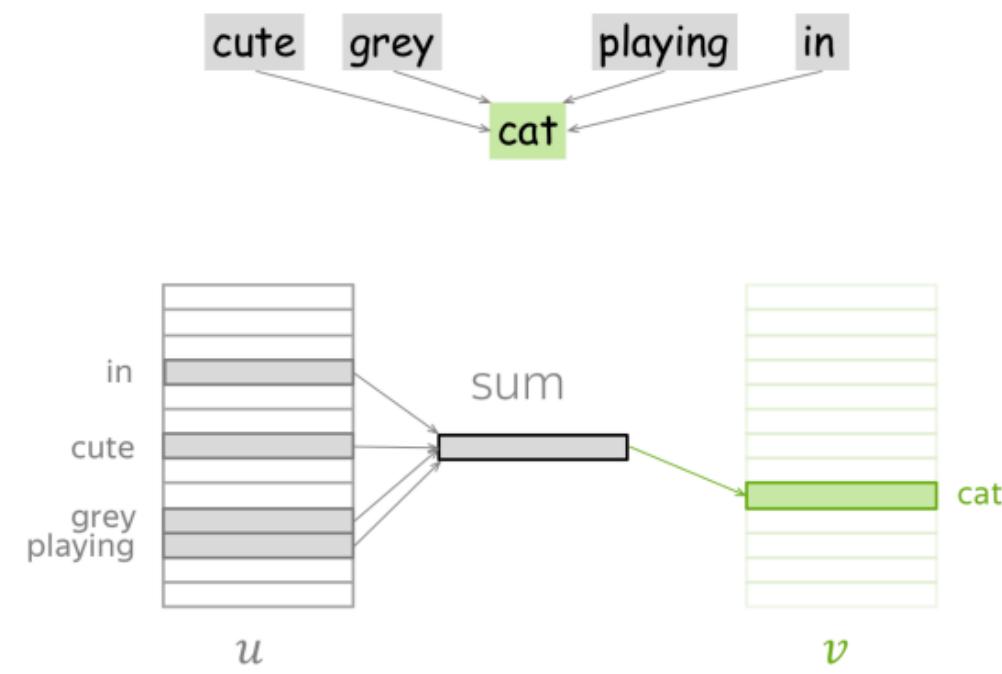


## Skip-gram

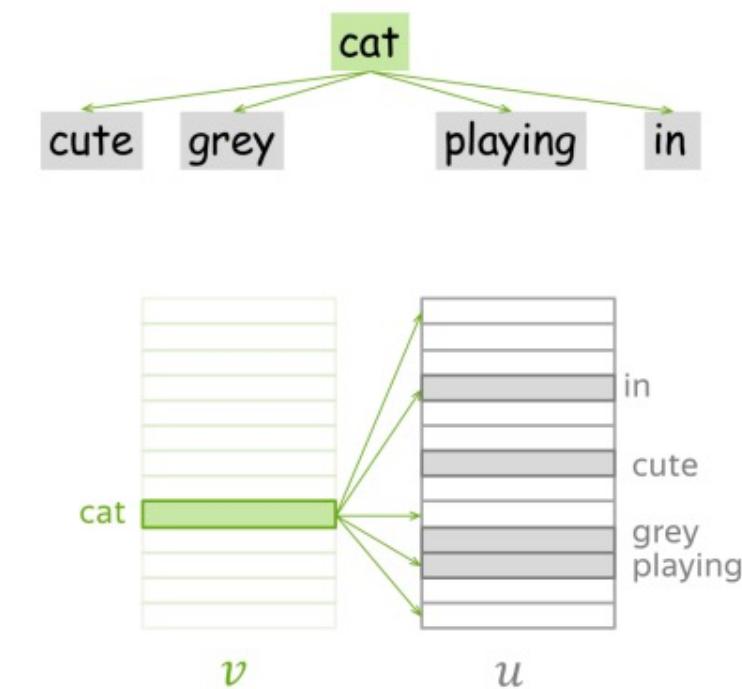


# Word2Vec: Two methods

## Continuous BOW (CBOW)



## Skip-gram



- > **From sum of context predict central**
- > Predicting one word each time
- > Relatively fast

- > **From central predict context (one at a time)**
- > Much slower
- > Better with infrequent words

# Word2Vec: Additional

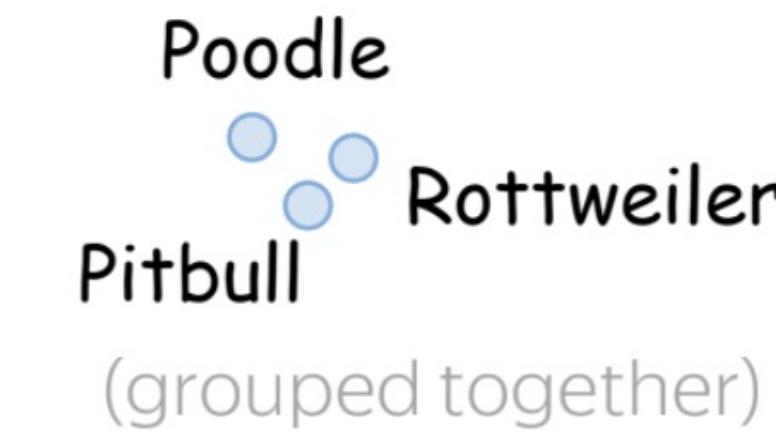
---

- **Dynamic window** — random selection of context size at each iteration or 5-10
- **Most popular:** Skip-Gram with negative sampling
- **Number of negative examples:**
  - for smaller datasets 15-20
  - for huge datasets it can be 2-5
- **Embedding dimensionality:** frequently used value is 300, but other variants (e.g., 100 or 50) are also possible

# Word2Vec: Practical tips

---

- Larger windows – more topical similarities
- Smaller windows – more functional and syntactic similarities



## Count-based

Information  
comes from:

global corpus  
statistics

Vectors are:

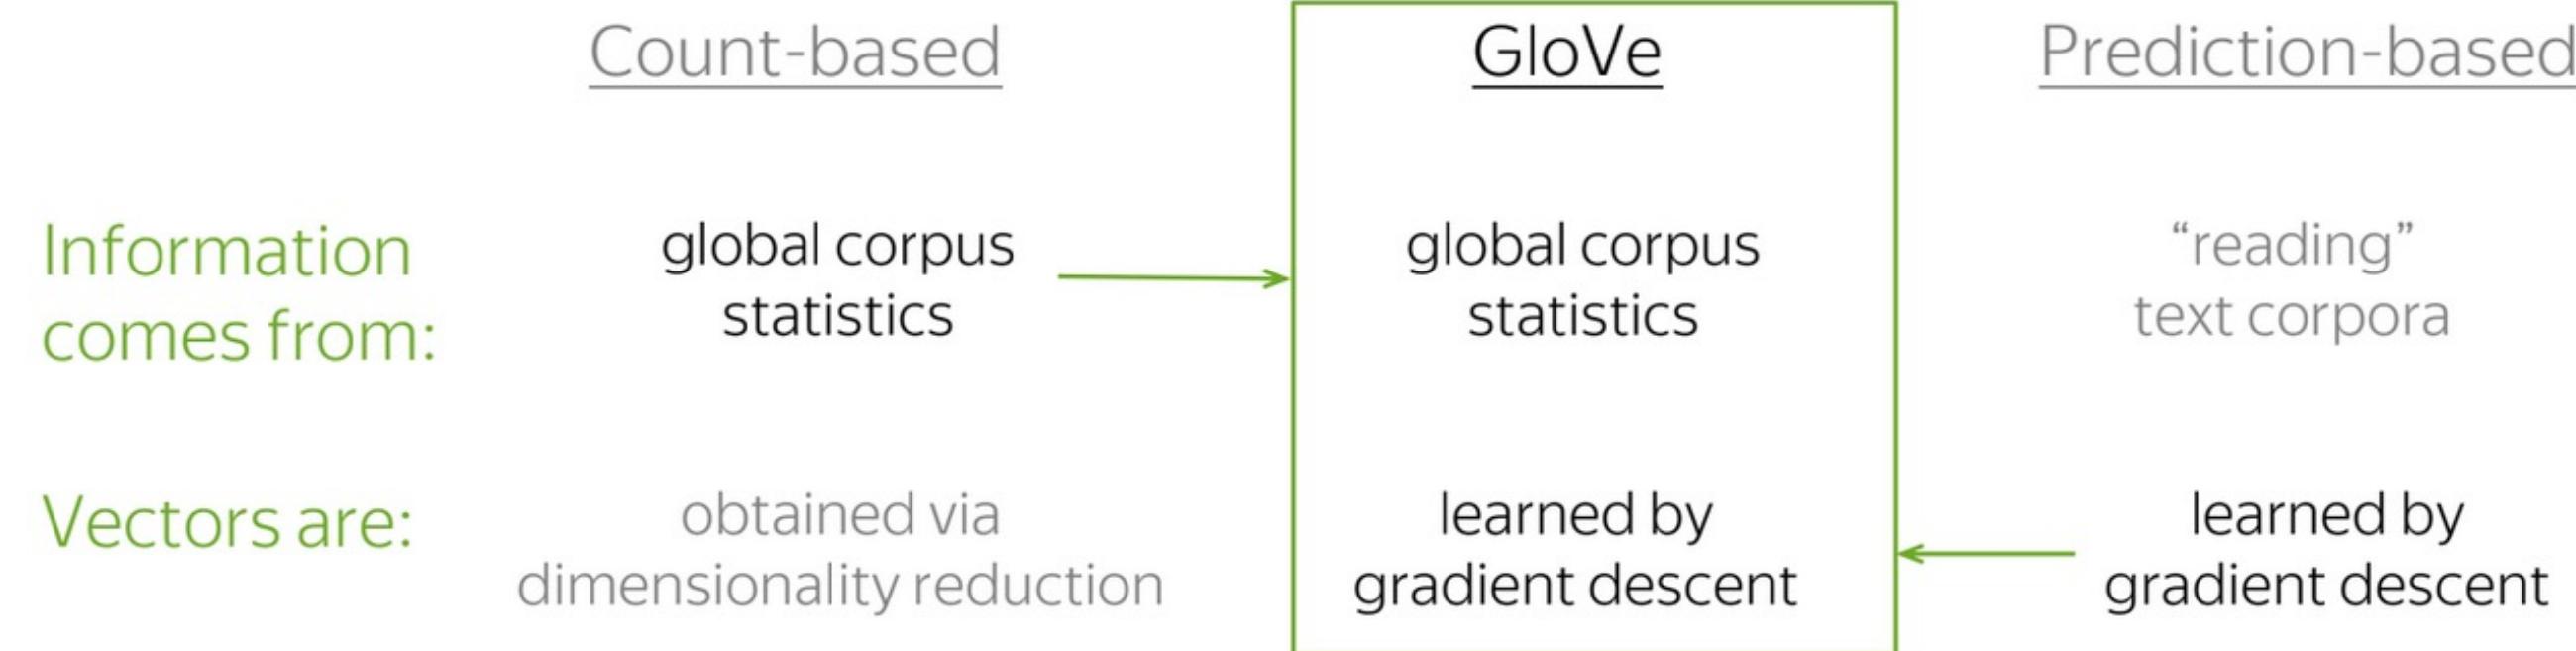
obtained via  
dimensionality reduction

## Prediction-based

“reading”  
text corpora

learned by  
gradient descent

## Global Vectors for word representations



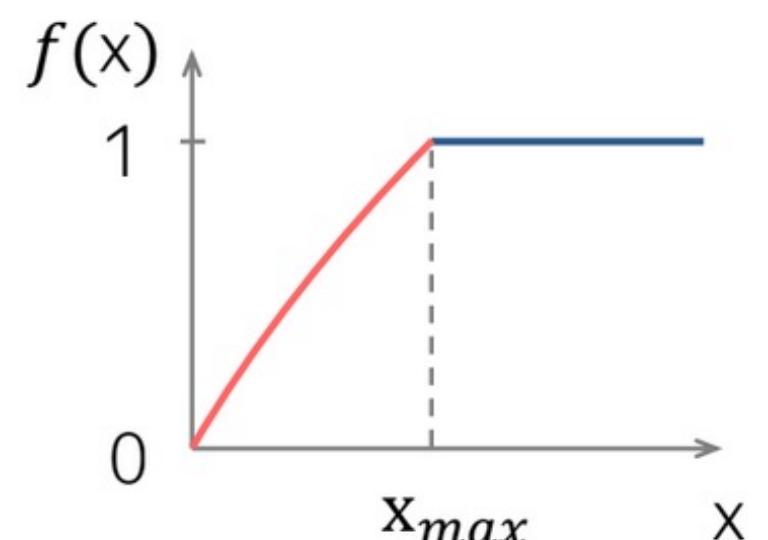
# GloVe

---

> Loss function:

$$J(\theta) = \sum_{w,c \in V} f(N(w, c)) \cdot (u_c^T v_w + b_c + \bar{b}_w - \log N(w, c))^2$$

context vector      word vector      bias terms  
(also learned)



$f(x) = \begin{cases} (x/x_{max})^\alpha & \text{if } x < x_{max}, \\ 1 & \text{otherwise.} \end{cases}$

$\alpha = 0.75, x_{max} = 100$

> Popular, but in practice usually worse than Word2vec

# FastText

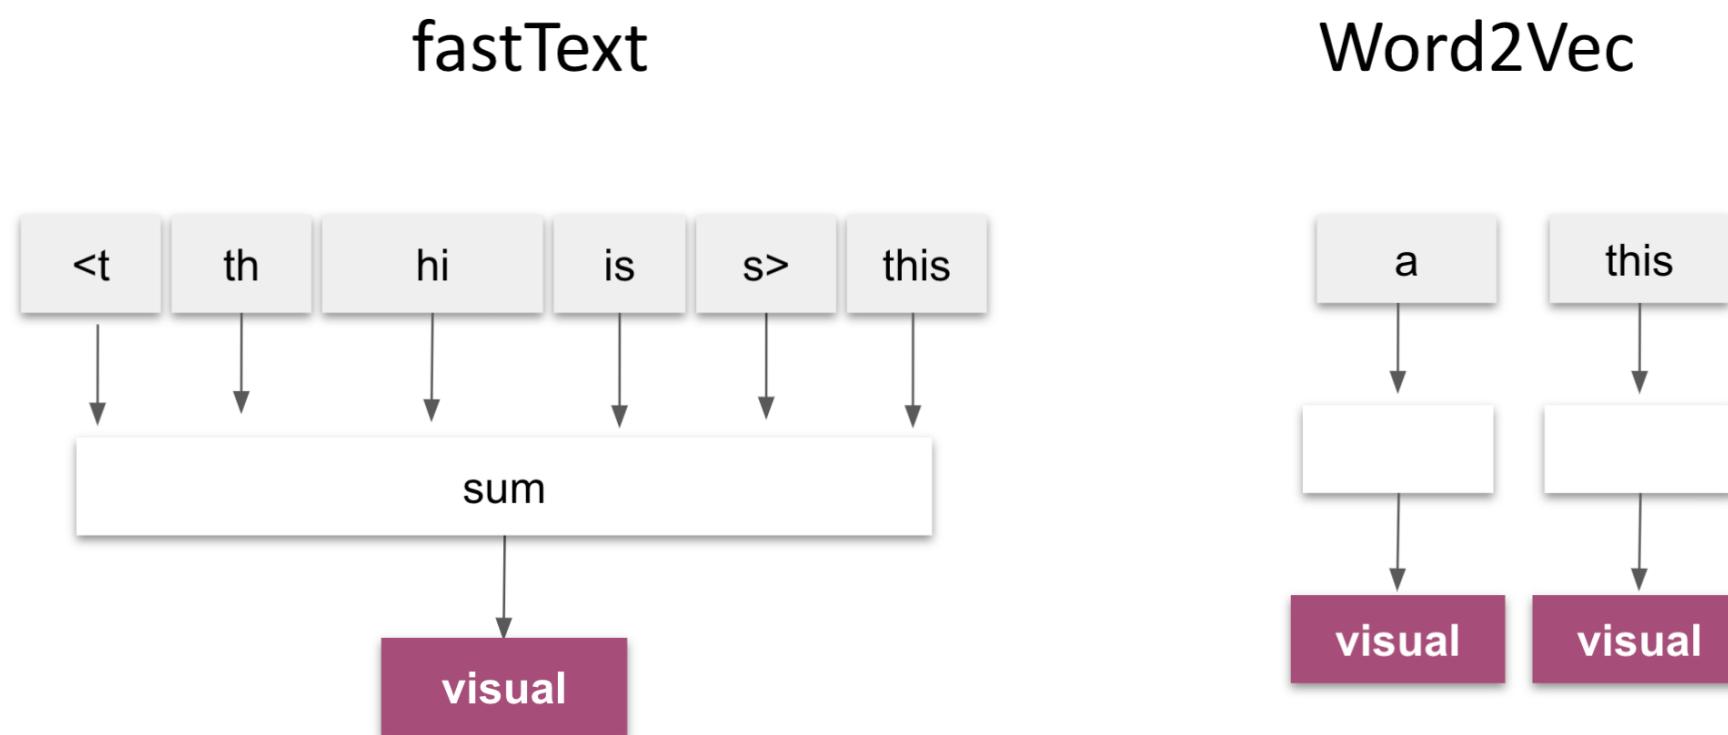
---

- Breaks words into subword components (n-grams) and learns embeddings for these subwords

# FastText

---

- Breaks words into subword components (n-grams) and learns embeddings for these subwords
- Can represent the word that was not presented in training set using the existing subword embeddings for its subwords



# Evaluation

---

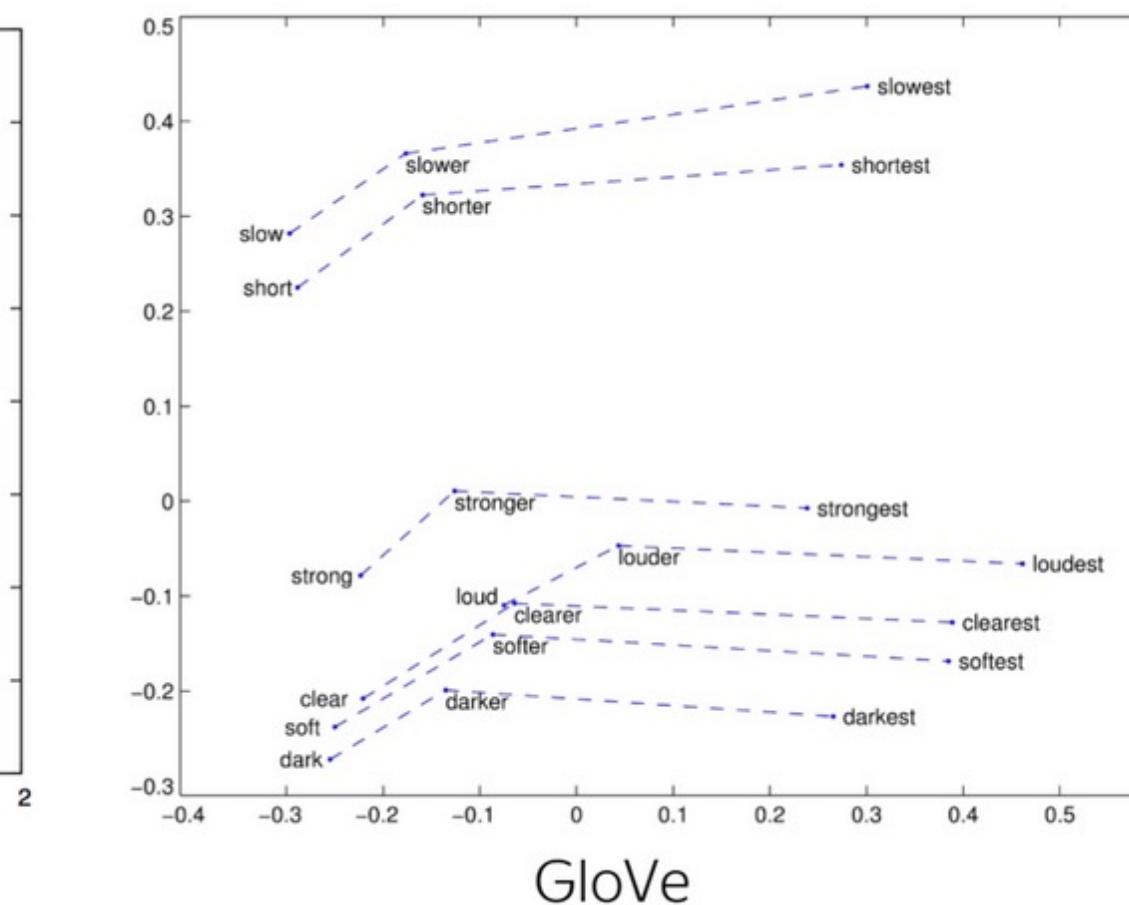
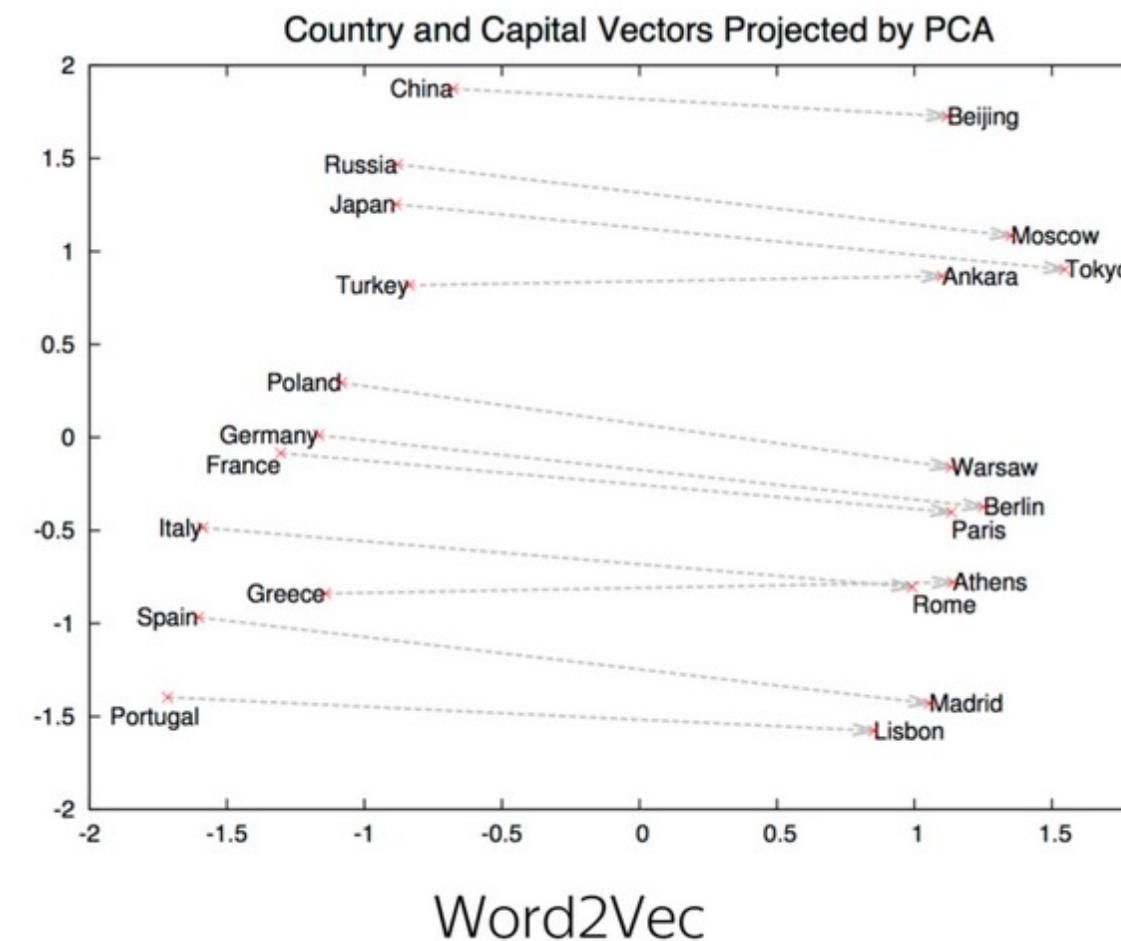
- **Which representations should be considered good?**

- Words that are close in meaning correspond to vectors that are close in distance
- Small dimension
- Interpreted arithmetic operations in vector space
- Quality of the solution to the finite problem.

# Evaluation

- Which representations should be considered good?

- > Words that are close in meaning correspond to vectors that are close in distance
- > Small dimension
- > Interpreted arithmetic operations in vector space
- > Quality of the solution to the finite problem.



# Evaluation

---

- **Which representations should be considered good?**

- Words that are close in meaning correspond to vectors that are close in distance
- Small dimension
- Interpreted arithmetic operations in vector space
- Quality of the solution to the finite problem.

- **How to use word embeddings?**

- Solve tasks involving searching for similar words, synonyms, etc.
- Obtain a representation of a document/sentence that can be used to solve a machine learning task
- Use the word representation as a fixed representation in a complex architecture (e.g., a recurrent network)
- Use to initialize representations in a complex architecture

# Current state for embeddings

---

- **Lexical matching**
  - > TF-IDF, BM-25 [fast, interpretable]
- **Semantic similarity**
  - > Sentence Transformers (bi-encoder), BERT-like models
- **Reranking**
  - > Cross-encoder [slower, higher quality]
- **Generation & reasoning**
  - > LLMs [prompting / fine-tuning]

# Recap

---

- Sparse / count-based representations
- Predictive embeddings
- Extensions & alternatives as FastText, GloVe for OOV and speed
- Representation choice = trade-off between interpretability, compute, and semantic power

Next:

- Language modelling & Attention